

# SMART-SBERT with SimCSE: A Robust Defense Against Catastrophic Forgetting in BERT



## Submitted by:

Rooshan Khan	2021-EE-067
Hussnain Amjad	2021-EE-063
Abdul Samad	2021-EE-191
Areesha Noor	2021-EE-103

**Supervised by:** Dr. Irfan Ullah Chaudhry

Department of Electrical Engineering  
**University of Engineering and Technology Lahore**

# **SMART-SBERT with SimCSE: A Robust Defense Against Catastrophic Forgetting in BERT**

Submitted to the faculty of Electrical Engineering  
of the University of Engineering and Technology Lahore  
in partial fulfillment of the requirements for the Degree of  
Bachelor of Science  
in  
**Electrical Engineering.**

---

Internal Examiner

---

External Examiner

---

Director  
Undergraduate Studies

Department of Electrical Engineering  
**University of Engineering and Technology Lahore**

# Declaration

We declare that the work contained in this thesis is our own, except where explicitly stated otherwise. In addition this work has not been submitted to obtain another degree or professional qualification.

Signed: \_\_\_\_\_

Date: \_\_\_\_\_

# Acknowledgments

We would like to express our heartfelt gratitude to our supervisor, Dr. Irfan Ullah Chaudhry, for his invaluable guidance, continuous support, and encouragement throughout the course of this project.

We are also deeply thankful to Dr. Haroon Babri for his insightful suggestions and constructive feedback, which significantly enhanced our understanding and progress.

Lastly, we extend our sincere appreciation to Mr. Asad Ullah Khan for his consistent assistance and helpful input during the course of this work.

*Dedicated to our parents and teachers, whose unwavering support,  
guidance, and encouragement have been the foundation of our  
journey and success.*

# Contribution to Sustainable Development Goals

We believe our thesis, “**SMART-SBERT with SimCSE: A Robust Defense Against Catastrophic Forgetting in BERT,**” makes an important contribution to the Sustainable Development Goals (SDGs) by advancing Artificial Intelligence (AI) and Natural Language Processing (NLP).

- **Quality Education**

Our work helps improve the reliability and flexibility of NLP models like BERT. These better models can power smarter educational tools. For example, by improving sentiment analysis, educational platforms can better understand how students feel and respond to their feedback. Similarly, enhancements in semantic textual similarity can support more accurate automated grading and personalized learning experiences. Ultimately, these advances can help improve learning outcomes and make quality education more accessible.

- **Industry, Innovation and Infrastructure**

At the heart of our thesis is innovation in AI and NLP. We’ve developed models that reduce catastrophic forgetting while performing well on multiple tasks, including sentiment classification, paraphrase detection, and semantic textual similarity. This helps build more robust and efficient digital systems. Because our models can adapt to new tasks sequentially without growing too large, they are well-suited for real-world industrial applications. We believe these innovations can foster technological progress and economic growth by enabling smarter data processing, automation, and intelligent systems across various industries.

# Contents

<b>Acknowledgments</b>	<b>iii</b>
<b>Contribution to Sustainable Development Goals</b>	<b>v</b>
<b>List of Figures</b>	<b>ix</b>
<b>List of Tables</b>	<b>x</b>
<b>Abbreviations</b>	<b>xi</b>
<b>Abstract</b>	<b>xii</b>
<b>1 Motivations and Problem Statement (Rooshan Khan)</b>	<b>1</b>
1.1 Proposed Approach . . . . .	2
1.2 Project Description . . . . .	3
1.2.1 Data . . . . .	3
1.2.2 Baseline . . . . .	3
1.2.3 Evaluation . . . . .	3
<b>2 Bidirectional Encoder Representations from Transformers: BERT (Abdul Samad)</b>	<b>4</b>
2.1 Introduction . . . . .	4
2.1.1 Bidirectional Encoder Representations from Transformers: BERT .	4
2.1.1.1 Introduction . . . . .	4
2.2 Architecture . . . . .	4
2.3 Tokenization . . . . .	5
2.4 Embedding Layer . . . . .	6
2.5 Transformer Encoder Layer . . . . .	6
2.6 Pre-training and Fine-tuning BERT Model . . . . .	7
2.7 Fine-tuning for Downstream Tasks . . . . .	8
<b>3 Theoretical Background</b>	<b>9</b>
3.1 SMART: Smoothness Inducing Adversarial Regularization and Bregman Proximal Point Optimization: Rooshan Khan . . . . .	9
3.1.1 Framework Overview . . . . .	9
3.1.2 Implementation of SMART . . . . .	11
3.2 Mathematical Foundations and Intuition Behind Supervised SimCSE Contrastive Loss: Abdul Samad . . . . .	11
3.2.1 Introduction . . . . .	11

3.2.2	Why Contrastive Learning?	11
3.2.3	Sentence Embeddings and Similarity	12
3.2.4	Supervised SimCSE Using SNLI Triplets	12
3.2.4.1	Dataset Structure	12
3.2.4.2	Goal	12
3.2.4.3	Mathematical Loss Derivation	12
3.2.5	Temperature Parameter ( $\tau$ )	13
3.2.6	Key Takeaways	13
3.3	Implementation of Sentence-BERT architecture for Semantic Textual Similarity and Paraphrase Detection: Hussnain Amjad	13
3.3.1	Background and Motivation	14
3.3.2	Architecture and Embedding Strategy	14
3.3.3	Paraphrase Detection (Classification Task)	14
3.3.4	Semantic Textual Similarity (Regression Task)	15
3.3.5	Training Details	17
<b>4</b>	<b>Experimentation (Rooshan Khan)</b>	<b>18</b>
4.1	Phase 1	18
4.2	Phase 2	18
4.2.1	Individually fine-tuned BERT models	20
4.2.2	Accuracies and Pearson Correlation for Fine-Tuning the same model on Multiple Tasks	22
4.2.2.1	Sequential Fine-Tuning: SBERT Only	22
4.2.2.2	Sequential Fine-Tuning: SBERT+SMART	25
4.2.2.3	Sequential Fine-Tuning: SBERT+SimCSE+SMART	28
4.3	Analysis of the Results	28
<b>5</b>	<b>Possible Improvements (Areesha Noor)</b>	<b>30</b>
<b>A</b>	<b>Implementation of Techniques in Python (Rooshan Khan)</b>	<b>32</b>
A.1	SMART	32
A.1.1	Implementation of methods for predictions	41
<b>B</b>	<b>Implementation of Supervised SimCSE in Python (Abdul Samad)</b>	<b>47</b>
B.1	Supervised SimCSE	47
B.1.1	SNLI Dataset Preprocessing	47
B.1.2	Triplet Contrastive Loss for SimCSE	48
B.1.3	Supervised SimCSE Loss Implementation	48
B.1.4	Training Loop for Supervised SimCSE	50
B.1.5	SNLI Dataset Preparation	51
<b>C</b>	<b>SBERT Implementation for Paraphrase Detection and Semantic Similarity (Hussnain Amjad)</b>	<b>55</b>
C.1	Overview	55





# List of Figures

1.1	Proposed Model . . . . .	3
2.1	BERT Architecture. Image sourced from <a href="https://wikidocs.net/book/1">https://wikidocs.net/book/1</a> . . . . .	5
3.1	SBERT Implementation Architecture for Paraphrase Detection. Referenced from <i>Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks</i> by Reimers and Gurevych [9]. . . . .	15
3.2	SBERT Implementation Architecture for Semantic Textual Similarity. Referenced from <i>Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks</i> by Reimers and Gurevych [9]. . . . .	16
4.1	Proposed Model . . . . .	19
4.2	Final Model Architecture . . . . .	20
4.3	SST train vs dev accuracy over 10 epochs . . . . .	21
4.4	STS train vs dev correlation over 10 epochs . . . . .	21
4.5	Train vs Dev Accuracy for Paraphrase Detection using SBERT Only (3 epochs) . . . . .	23
4.6	Train vs Dev Correlation for Semantic Textual Similarity using SBERT Only (10 epochs) . . . . .	23
4.7	Train vs Dev Accuracy for Sentiment Analysis using SBERT Only (10 epochs) . . . . .	24
4.8	Train vs Dev Pearson Correlation for STS using SMART SBERT (10 epochs) . . . . .	25
4.9	Train vs Dev Accuracy for SST using SMART SBERT (10 epochs) . . . . .	26
4.10	Updated Training Procedure . . . . .	27

# List of Tables

1.1	Tasks and Ratings for simple BERT model without finetuning . . . . .	2
1.2	Dataset splits used for training, validation, and testing. . . . .	3
1.3	Evaluation metrics for the tasks . . . . .	3
2.1	WordPiece Tokenization . . . . .	6
4.1	Quantitative Result of Phase 1. Accuracy comparison of BERT on SST and CFIMDB. . . . .	18
4.2	Best performance of separate models fine-tuned individually for each task	22
4.3	Development set performance before fine-tuning . . . . .	22
4.4	Development set performance after fine-tuning on Paraphrase Detection for 3 epochs . . . . .	23
4.5	Development set performance after fine-tuning on the STS task . . . . .	24
4.6	Development set performance after fine-tuning for 10 epochs on SST . . .	24
4.7	Development set performance after fine-tuning on Paraphrase Detection for 1 epoch . . . . .	25
4.8	Development set performance after fine-tuning for 10 epochs on STS . . .	26
4.9	Development set performance after fine-tuning for 10 epochs on SST . . .	26
4.10	Evaluation metrics for each task (SMART+SBERT) . . . . .	27
4.11	Evaluation metrics for each task (SBERT Only) . . . . .	27
4.12	Evaluation metrics for each task . . . . .	28
4.13	Evaluation metrics for each task . . . . .	28
4.14	State-of-the-art BERT-base results for each task . . . . .	29

# Abbreviations

<b>NLP</b>	Natural <b>L</b> anguage <b>P</b> rocessing
<b>BERT</b>	<b>B</b> idirectional <b>E</b> ncoder <b>R</b> epresentations from <b>T</b> ransformers
<b>SMART</b>	<b>S</b> Moothness-inducing <b>A</b> dversarial <b>R</b> egularization and <b>B</b> Regman p <b>R</b> oximal point op <b>T</b> imization
<b>SBERT</b>	Sentence- <b>BERT</b>
<b>SimCSE</b>	<b>S</b> imple <b>C</b> ontrastive <b>S</b> entence <b>E</b> mbeddings
<b>MLM</b>	Masked <b>L</b> anguage <b>M</b> odel
<b>NSP</b>	Next <b>S</b> entence <b>P</b> rediction
<b>SQuAD</b>	Stanford <b>Q</b> uestion <b>A</b> nsWERing <b>D</b> ataset
<b>STS</b>	Semantic <b>T</b> extual <b>S</b> imilarity
<b>STS-B</b>	Semantic <b>T</b> extual <b>S</b> imilarity <b>B</b> enchmark
<b>QQP</b>	<b>Q</b> uora <b>Q</b> uestion <b>P</b> airs
<b>SST</b>	Stanford <b>S</b> entiment <b>T</b> reebank
<b>SNLI</b>	Stanford Natural <b>L</b> anguage <b>I</b> nterference
<b>FFN</b>	<b>F</b> eed- <b>F</b> orward <b>N</b> etwork
<b>GeLU</b>	<b>G</b> aussian <b>E</b> rror <b>L</b> inear <b>U</b> nit
<b>ReLU</b>	<b>R</b> ectified <b>L</b> inear <b>U</b> nit
<b>KL</b>	<b>K</b> ullback- <b>L</b> eibler
<b>VBPP</b>	<b>V</b> anilla <b>B</b> regman <b>P</b> roximal <b>P</b> oint
<b>MBPP</b>	<b>M</b> omentum <b>B</b> regman <b>P</b> roximal <b>P</b> oint
<b>SS</b>	<b>S</b> MART <b>S</b> BERT
<b>SSS</b>	<b>S</b> MART <b>S</b> BERT with <b>S</b> imCSE

# Abstract

We aim to build a robust model based on the BERT-base architecture that performs well on three NLP tasks—sentiment classification (SST-5), paraphrase detection (QQP), and semantic textual similarity (STS-B)—by training sequentially across tasks. A core challenge in this setup is *catastrophic forgetting*, where learning a new task can degrade performance on previously learned tasks due to destructive gradient updates.

To mitigate this, we implemented the **SMART algorithm**, which introduces adversarial regularization to preserve previously acquired knowledge and avoid aggressive parameter updates. This component was implemented by Rooshan Khan. Abdul Samad and Areesha contributed contrastive learning using supervised and unsupervised loss respectively. However, we used the best among these two that is supervised SimCSE loss to maximize cosine similarity between semantically similar sentence pairs and minimize it for dissimilar ones. Hussnain Amjad implemented the SBERT architecture, for task-specific classification and regression heads for QQP and STS-B, respectively.

We got two best models:

- **SMART SBERT (SS)**
- **SMART SBERT with SimCSE (SSS)**

**SS** performed better on the SST-5 task, while **SSS** showed superior results on the STS-B task due to the SimCSE loss function being optimized for similarity learning. Both models performed equally well on paraphrase detection. Results are shown below:

Task	Metric	Score (SS)
Sentiment Classification	Accuracy	0.537
	F1 Score	0.528
Paraphrase Detection	Accuracy	0.864
	F1 Score	0.864
Semantic Textual Similarity	Pearson Correlation	0.819
<b>Overall Performance</b>		<b>0.770</b>

$$\text{Overall performance (SS)} = \frac{\left(\frac{0.819+1}{2} + 0.537 + 0.864\right)}{3} = 0.770$$

Task	Metric	Score (SSS)
Sentiment Classification	Accuracy	0.507
	F1 Score	0.494
Paraphrase Detection	Accuracy	0.864
	F1 Score	0.863
Semantic Textual Similarity	Pearson Correlation	0.843
<b>Overall Performance</b>		<b>0.764</b>

$$\text{Overall performance (SSS)} = \frac{\left(\frac{0.843+1}{2} + 0.507 + 0.864\right)}{3} = 0.764$$

Depending on task priority, either SS or SSS may be selected. Our results show that a combination of SMART regularization and contrastive learning can effectively reduce catastrophic forgetting and achieve strong performance across all three tasks in a sequential training setup.

# Chapter 1

## Motivations and Problem Statement (Rooshan Khan)

Pre-trained models like BERT (Bidirectional Encoder Representations from Transformers), presented by Devlin and colleagues in their work *BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding* [1], have transformed natural language processing by setting new benchmarks across numerous downstream tasks. However, despite its success, the BERT model is known to suffer from issues related to **robustness**—its ability to maintain consistent performance when exposed to challenging, noisy, or perturbed inputs. Furthermore, typical fine-tuning strategies often lead to large, unstable updates, making the model prone to **overfitting**, especially when training on limited data.

In our work, we aim to improve the robustness and generalizability of the BERT model. Unlike traditional approaches where separate BERT models are fine-tuned independently for each downstream task, we adopt a **sequential fine-tuning strategy** in which a single BERT model is progressively fine-tuned across multiple tasks. Specifically, the same BERT model is trained first on paraphrase detection (QQP), then on semantic textual similarity (STS-B), and finally on sentiment analysis (SST-5). This approach allows for the transfer of useful knowledge between tasks while keeping the model size compact.

We propose a robust fine-tuning framework that integrates the following components:

- **Smoothness-Inducing Adversarial Regularization:** Enhances model robustness by training it to resist small perturbations in the input space.[3]
- **Bregman Proximal Point Optimization:** Prevents large and unstable parameter updates during fine-tuning, ensuring smoother convergence and reducing overfitting.[3]

- **Contrastive Learning:** Enhances sentence representations by pulling embeddings of semantically related sentence pairs closer together, while pushing apart those of unrelated pairs within the embedding space.

We also adopt the **Siamese Sentence-BERT (SBERT)** architecture to create meaningful sentence embeddings that capture semantic similarity, enabling efficient comparison of sentence pairs.[9]

Our experiments use the **BERT-Base model** due to its manageable size. However, the proposed techniques can be extended to larger BERT variants, such as BERT-Large, in future work.

Table 1.1 presents the performance of the baseline BERT model on these tasks.

Tasks	Accuracies(1st 2) and Pearson Correlation (last)
Sentiment Analysis	0.144
Paraphrase Detection	0.369
Semantic Textual Similarity	-0.009

TABLE 1.1: Tasks and Ratings for simple BERT model without finetuning

**Motivation:** Looking at Table 1.1, we recognized the need for fine-tuning. We wanted to investigate how fine-tuning the same model for different tasks affects its performance on those tasks—specifically, whether it improves performance on individual tasks or has the opposite effect. We explored this question using our selected tasks. Our experiments showed a slight improvement in accuracy, but it was not significant. Hence, fine-tuning alone is not sufficient. It needs to be combined with additional techniques from research papers, which we discuss in the following sections.

## 1.1 Proposed Approach

We won't be changing the architecture of BERT. This means that the first 12 layers will be the same. We will add the task specific heads for the three downstream tasks. These task specific layers will vary as we proceed with project. The three tasks are Sentiment Analysis (Classification), Paraphrase Detection (Classification), and Semantic Textual Similarity (Regression). We will take the final sentence embedding from BERT and use it as input to the task specific layers. We will be finetuning the model for each task with its dataset for a specific number of epochs. We will do this for all the three tasks. The BERT Layers will be updated regardless of the task because gradients from all tasks will backpropagate through BERT. BERT is shared for the all the tasks. The task specific heads- Sentiment Analysis (Classification head), Paraphrase Detection (Classification head), and Semantic Textual Similarity (Regression head) - will only be updated for the task for which finetuning is being done. Other task heads won't be updated. This can be understood from Figure 1.1. We know this is not enough to get good accuracies so we plan to implement three techniques to get a fine Multi-task Model. These techniques are SMART, Contrastive Learning and Siamese Sentence BERT.



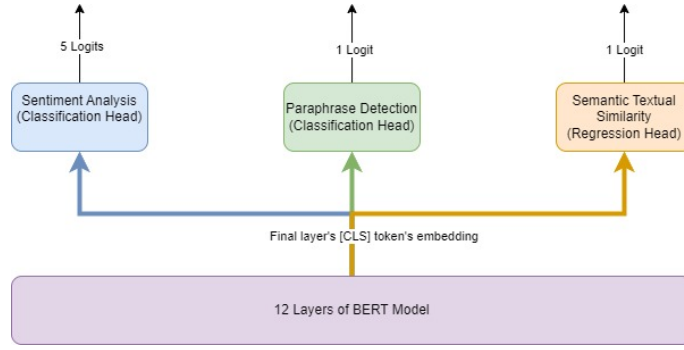


FIGURE 1.1: Proposed Model

## 1.2 Project Description

### 1.2.1 Data

Dataset	Train	Dev	Test
Stanford Sentiment Treebank (SST)	8,544	1,101	2,210
Quora Question Pairs	283,010	40,429	80,859
SemEval STS Benchmark	6,040	863	1,725
SNLI	550,152	Not used	Not used

TABLE 1.2: Dataset splits used for training, validation, and testing.

### 1.2.2 Baseline

We define the following baseline to estimate the upper-bound performance achievable with task-specific optimization:

- **Individually fine-tuned BERT models:** Train separate BERT models, each optimized for a single task. This setup serves as a task-specific upper-bound reference using the techniques employed.

### 1.2.3 Evaluation

We evaluate our approach on three benchmark NLP tasks:

Task	Evaluation Metric
<b>Paraphrase Detection (QQP)</b>	Accuracy
<b>Semantic Textual Similarity (STS-B)</b>	Pearson correlation coefficient
<b>Sentiment Analysis (SST-5)</b>	Accuracy

TABLE 1.3: Evaluation metrics for the tasks

## Chapter 2

# Bidirectional Encoder Representations from Transformers: BERT (Abdul Samad)

### 2.1 Introduction

#### 2.1.1 Bidirectional Encoder Representations from Transformers: BERT

##### 2.1.1.1 Introduction

Prior to BERT [1], word embedding models such as Word2Vec [5] and GloVe [7] provided static representations for words. For instance, a word like *bank* would always map to the same vector, regardless of the context in which it was used. While these models were effective for certain tasks, they struggled with capturing contextual meaning within sentences. The advent of the Transformer architecture, with its self-attention mechanism, not only enabled parallel processing—thus utilizing the full computational power of GPUs—but also facilitated the modeling of long-range dependencies. Introduced in 2018, BERT (Bidirectional Encoder Representations from Transformers) leveraged these advancements to create deeply bidirectional contextual embeddings, revolutionizing the field of Natural Language Processing (NLP).

### 2.2 Architecture

BERT's architecture is based on the Transformer encoder introduced by Vaswani et al. [11], which employs stacked self-attention layers to process input sequences in parallel. Unlike decoder-based models, BERT uses only the encoder component to generate bidirectional representations. Key architectural features include multi-head self-attention, positional embeddings, layer normalization, and residual connections. These components

collectively enable BERT to capture contextual relationships between words, incorporate positional information into token representations, and stabilize training in deep networks. BERT is primarily available in two configurations:

- **BERT-base:** 12 layers, 768 hidden dimensions, 12 attention heads
- **BERT-large:** 24 layers, 1024 hidden dimensions, 16 attention heads

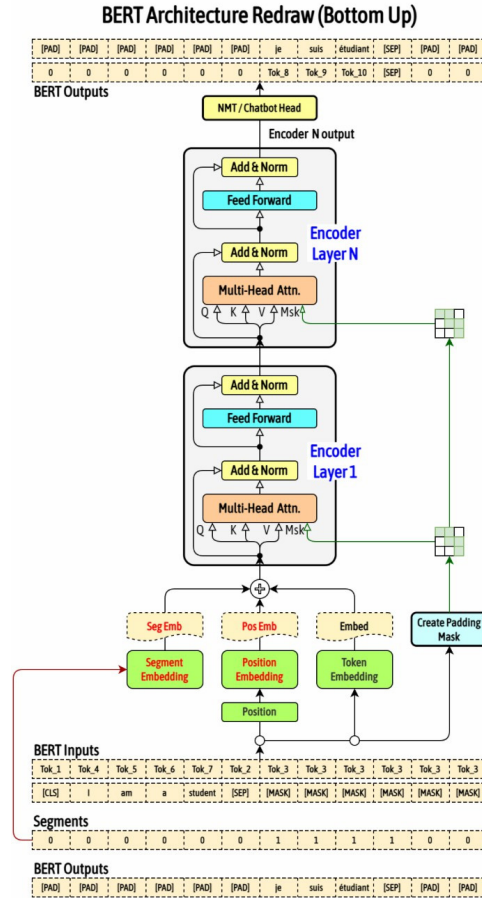


FIGURE 2.1: BERT Architecture. Image sourced from <https://wikidocs.net/book/1>.

## 2.3 Tokenization

BERT uses a WordPiece tokenizer to convert sentences into tokens. The main goals of tokenization are splitting words into subword units, handling out-of-vocabulary words, padding to ensure equal-length sequences, and adding special tokens [CLS] and [SEP]. As an example, the following words are converted into the following word pieces:

Word	WordPiece Tokens
running	[run, ##ing]
happiness	[happi, ##ness]
unreadable	[un, ##read, ##able]
football	[foot, ##ball]
unbelievable	[un, ##believ, ##able]

TABLE 2.1: WordPiece Tokenization

In addition to separating each sentences into word piece token, word token that have not been seen previously will be set as [UNK].

## 2.4 Embedding Layer

After tokenization, BERT converts each token into a high-dimensional vector using word embeddings. The embeddings are the sum of token embeddings, positional embeddings, and segmentation embeddings. The total embedding for each token is given as:

$$\text{embedding}(i) = \text{token\_embedding}(i) + \text{position\_embedding}(i) + \text{segmentation\_embedding}(i) \quad (2.1)$$

## 2.5 Transformer Encoder Layer

BERT is composed of 12 transformer encoder layers. Each layer integrates components such as multi-head self-attention, layer normalization, residual connections, and feed-forward networks (FFNs). The fundamental operation within the transformer is self-attention, which enables each token to consider the entire input sequence when generating its representation. In multi-head attention, the model employs  $n$  parallel attention heads, each performing scaled dot-product attention independently. The outputs from these heads are then concatenated and passed through a linear transformation using a weight matrix  $W_o$ :

$$\text{MultiHead}(h) = [\text{Attention}_1(h), \dots, \text{Attention}_n(h)]W_o \quad (2.2)$$

After the computation of attention, BERT adds the original input to the attention output using a residual connection:

$$h' = \text{LN}(h + \text{MultiHead}(h)) \quad (2.3)$$

$$\text{SA}(h) = \text{FFN}(h) \quad (2.4)$$

where LN is layer normalization. Each encoder layer contains a feed-forward network that performs two linear transformations with a non-linear activation function applied

in between:

$$\text{FFN}(h) = W_2 f(W_1 h + b_1) + b_2^1 \quad (2.5)$$

where  $f(\cdot)$  is a non-linearity, we use GeLU in BERT rather than ReLU. The Gaussian Error Linear Unit (GeLU) and Rectified Linear Unit (ReLU) are both activation functions that are widely used in neural networks, but their differences in handling negative values make GeLU particularly effective in complex models like BERT whereas ReLU, defined as  $\text{ReLU}(x) = \max(0, x)$ , outputs zero for all negative inputs, creating a strict boundary at zero. GeLU introduces a smooth, probabilistic transition for negative values near zero. This smoothness ensures GeLU avoids abrupt saturation, maintaining gradient flow even for inputs slightly below zero.

## 2.6 Pre-training and Fine-tuning BERT Model

One of the biggest challenges in language model training is to determine a prediction goal. A directional method that may restrict context learning is used by several models to predict the next word in a sequence. BERT uses two advanced methods of training to address this issue:

- **Masked Language Model (MLM)** The MLM task helps BERT understand words by using their surrounding context. During training, about 15% of the words in each sentence are selected for masking:
  - 80% of these words are replaced with the special [MASK] token.
  - 10% are replaced with random words.
  - 10% stay the same.

BERT is then trained to guess the original word at each masked position using the other words in the sentence. A classifier layer on top of BERT predicts the correct word from the vocabulary. Since it only tries to guess the masked words, BERT becomes very good at understanding context in both directions.

- **Next Sentence Prediction (NSP)** The NSP task helps BERT understand how sentences are related. In this task, BERT is given two sentences and must decide if the second sentence follows the first one in the original text.
  - 50% of the time, the second sentence is correct.
  - 50% of the time, it is a random sentence from somewhere else.

To prepare the input:

- A [CLS] token is added at the start.
- A [SEP] token separates the two sentences.

---

<sup>1</sup>FFN: Feed-Forward Network

- Segment embeddings mark which words belong to Sentence A or Sentence B.
- Positional embeddings encode the position of each token within the input sequence.

The output from the [CLS] token is passed through a small classifier that predicts whether the second sentence follows the first. A softmax function is used to give a probability.

$$\text{Total Loss} = \text{NSP Loss} + \text{MLM Loss} \quad (2.6)$$

## 2.7 Fine-tuning for Downstream Tasks

BERT’s versatility enables it to be fine-tuned for various NLP tasks with minimal architectural modifications. For instance, in text classification, a classification head is added on top of the [CLS] token’s representation. For named entity recognition (NER), BERT leverages token-level embeddings to make predictions. In question answering tasks like those based on the SQuAD dataset [8], BERT is fine-tuned by learning to identify the start and end indices of the answer span within the context.

### **What is the difference between pre-training and fine-tuning BERT?**

Pre-training involves training the model on large-scale unlabeled text using specific self-supervised objectives. In contrast, fine-tuning adapts the pre-trained BERT model to a specific task by updating all of its parameters using labeled data relevant to that task.

## Chapter 3

# Theoretical Background

In our research, we have studied several key papers in the field of Natural Language Processing (NLP) to gain insights into pretraining, fine-tuning strategies, and sentence embeddings. The techniques that we will discuss in this chapter have been taken from some important research papers.

### 3.1 SMART: Smoothness Inducing Adversarial Regularization and Bregman Proximal Point Optimization: Rooshan Khan

This framework was introduced in the research paper *SMART: Robust and Efficient Fine-Tuning for Pre-trained Natural Language Models through Principled Regularized Optimization* by Jiang et al. [3] Aggressive fine-tuning often causes the model to overfit the training data of downstream tasks, resulting in poor generalization to unseen data. To tackle this problem, the authors introduced a novel learning framework aimed at improving generalization performance.

#### 3.1.1 Framework Overview

The proposed framework consists of two fundamental techniques to enhance the fine-tuning process of pre-trained language models:

1. **Smoothness-Inducing Adversarial Regularization:**

This technique enforces local smoothness by penalizing the model’s sensitivity to small input perturbations. Specifically, it identifies the worst-case perturbation within an  $\epsilon$ -ball around each input. First, for each sample  $x_i$ , we find a nearby perturbed input  $\tilde{x}_i$  that maximizes the smoothness loss. Then, we minimize the combined task and smoothness loss over model parameters  $\theta$ . A symmetrized Kullback-Leibler divergence is applied for classification tasks, while a squared loss is used for regression. This regularization effectively controls the model complexity and mitigates overfitting. The mathematics of this technique can be understood

from the equations given below. We solve the following optimization for fine-tuning.

$$\min_{\theta} \mathcal{F}(\theta) = \mathcal{L}(\theta) + \lambda_s \mathcal{R}_s(\theta),$$

where  $\mathcal{L}(\theta)$  represents the loss function which is defined as

$$\mathcal{L}(\theta) = \frac{1}{n} \sum_{i=1}^n \ell(f(x_i; \theta), y_i),$$

and  $\ell(\cdot, \cdot)$  denotes the task-specific loss function (We used Binary Cross Entropy loss function for Paraphrase Detection, MSE loss function for STS, and Cross Entropy loss function for Sentiment Analysis),  $\lambda_s > 0$  is a tuning parameter, and  $\mathcal{R}_s(\theta)$  is the Smoothness-Inducing Adversarial Regularizer. Here we define  $\mathcal{R}_s(\theta)$  as

$$\mathcal{R}_s(\theta) = \frac{1}{n} \sum_{i=1}^n \max_{\|\tilde{x}_i - x_i\|_p \leq \epsilon} \ell_s(f(\tilde{x}_i; \theta), f(x_i; \theta)),$$

where  $\epsilon > 0$  is a tuning parameter. Note that for classification tasks,  $f(\cdot; \theta)$  outputs a probability simplex<sup>1</sup> and  $\ell_s$  is chosen as the symmetrized KL-divergence, i.e.,

$$\ell_s(P, Q) = \mathcal{D}_{\text{KL}}(P \| Q) + \mathcal{D}_{\text{KL}}(Q \| P);$$

KL Divergence has the following formula:

$$D_{\text{KL}}(P \| Q) = \sum_{x \in \mathcal{X}} P(x) \log \left( \frac{P(x)}{Q(x)} \right)$$

For regression tasks,  $f(\cdot; \theta)$  outputs a scalar and  $\ell_s$  is chosen as the squared loss, i.e.,

$$\ell_s(p, q) = (p - q)^2.$$

## 2. Bregman Proximal Point Optimization:

To stabilize fine-tuning, this method adds a penalty based on a Bregman divergence that constrains parameter updates. It ensures that each update remains close to the previous parameters, thereby preventing excessively aggressive changes. A momentum variant further accelerates convergence and enhances stability. The equations related to this technique are given below. At the  $(t+1)$ -th iteration, the Vanilla Bregman Proximal Point (VBPP) method takes

$$\theta_{t+1} = \arg \min_{\theta} \mathcal{F}(\theta) + \mu D_{\text{Breg}}(\theta, \theta_t)$$

---

<sup>1</sup>For  $n$  possible outcomes, a point in the probability simplex is a vector  $\mathbf{p} = (p_1, p_2, \dots, p_n)$ , where  $p_i \geq 0$  for all  $i$ , and  $\sum_{i=1}^n p_i = 1$ .



where  $\mu > 0$  and the Bregman Divergence is given by

$$D_{\text{Breg}}(\theta, \theta_t) = \frac{1}{n} \sum_{i=1}^n \ell(f(\theta; x_i), f(\theta_t; x_i))$$

In our implementation we have acceleration by momentum as well so our parameters get updated like this:

$$\theta_{t+1} = \arg \min_{\theta} \mathcal{F}(\theta) + \mu D_{\text{Breg}}(\theta, \tilde{\theta}_t)$$

where

$$\tilde{\theta}_t = (1 - \beta)\theta_t + \beta\theta_{t-1}$$

is the exponential moving average and  $\beta \in (0, 1)$  is the momentum parameter. This is the momentum Bregman proximal point (MBPP) method.

### 3.1.2 Implementation of SMART

We implemented the SMART algorithm on all three tasks. We observed improvements in the accuracies of all the tasks (Sentiment Analysis, Semantic Textual Similarity, and Paraphrase Detection). We also observed that the maximum accuracy is now achieved in fewer epochs. However, the drawbacks are that the time required to train each epoch has increased by 4 to 5 times compared to training without SMART, and additional memory is required. Specifically, for the Paraphrase Detection task, each epoch now takes approximately 5 hours with SMART, compared to 50 minutes without SMART.

## 3.2 Mathematical Foundations and Intuition Behind Supervised SimCSE Contrastive Loss: Abdul Samad

### 3.2.1 Introduction

SimCSE (Simple Contrastive Learning of Sentence Embeddings) is a powerful technique for learning sentence representations that capture semantic meaning. This technique was introduced in the research paper *SimCSE: Simple Contrastive Learning of Sentence Embeddings* by Gao et al. [2]. This document explains in depth the mathematics and intuition behind the Supervised SimCSE approach and its adaptation to paraphrase detection, especially in the context of the SNLI dataset. The goal is to enable anyone with basic machine learning understanding to grasp the core concepts.

### 3.2.2 Why Contrastive Learning?

In traditional classification, we teach models to map inputs to fixed class labels. In contrastive learning, we teach models to understand relationships between inputs, especially:

- Which examples should be close (positive pairs)

- Which examples should be far apart (negative pairs)

We do this by pulling similar sentence embeddings together and pushing dissimilar ones apart. The result is a semantic embedding space where cosine similarity reflects actual meaning.

### 3.2.3 Sentence Embeddings and Similarity

Given a sentence, BERT produces a high-dimensional vector (embedding). To compare two sentences:

1. Convert both into vectors (via BERT and pooling)
2. Normalize the vectors
3. Measure cosine similarity:

$$\text{sim}(\mathbf{u}, \mathbf{v}) = \frac{\mathbf{u}^\top \mathbf{v}}{\|\mathbf{u}\| \cdot \|\mathbf{v}\|}$$

If vectors are unit-normalized:

$$\text{sim}(\mathbf{u}, \mathbf{v}) = \mathbf{u}^\top \mathbf{v}$$

### 3.2.4 Supervised SimCSE Using SNLI Triplets

#### 3.2.4.1 Dataset Structure

Each training example is a triplet:

- Anchor ( $a$ ): Premise
- Positive ( $p$ ): Hypothesis with entailment
- Negative ( $n$ ): Hypothesis with contradiction

#### 3.2.4.2 Goal

We want:

- The similarity between  $a$  and  $p$  to be high
- The similarity between  $a$  and  $n$  to be low

#### 3.2.4.3 Mathematical Loss Derivation

Let:

- $\mathbf{a}_i$ : embedding of anchor

- $\mathbf{p}_i$ : embedding of positive
- $\mathbf{n}_i$ : embedding of negative

All embeddings are L2-normalized:

$$s_i^+ = \mathbf{a}_i^\top \mathbf{p}_i, \quad s_i^- = \mathbf{a}_i^\top \mathbf{n}_i$$

Convert these into probabilities using softmax with temperature  $\tau$ :

$$P_i = \frac{\exp(s_i^+/\tau)}{\exp(s_i^+/\tau) + \exp(s_i^-/\tau)}$$

The contrastive loss per instance:

$$\mathcal{L}_i = -\log(P_i)$$

Overall loss per batch:

$$\mathcal{L} = \frac{1}{N} \sum_{i=1}^N \mathcal{L}_i$$

### 3.2.5 Temperature Parameter ( $\tau$ )

The temperature  $\tau$  controls the sharpness of separation:

- Small  $\tau$ : encourages hard separation, but may overfit
- Large  $\tau$ : softer distinction, but may underfit

Typical value:  $\tau = 0.05$

### 3.2.6 Key Takeaways

- SimCSE learns semantic sentence embeddings through contrastive learning.
- Cosine similarity and softmax are central to the formulation.
- SNLI provides well-structured triplets for training.

## 3.3 Implementation of Sentence-BERT architecture for Semantic Textual Similarity and Paraphrase Detection: Hussnain Amjad

The SBERT architecture was introduced in the research paper *Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks* by Reimers and Gurevych [9].

### 3.3.1 Background and Motivation

BERT provides deep contextual language representations but lacks the capacity to generate good sentence embeddings that can be efficiently compared. Sentence-BERT (SBERT) overcomes this limitation by employing a Siamese architecture that produces meaningful sentence vectors, enabling direct comparison through distance-based metrics. In our project, SBERT was implemented and fine-tuned for two tasks: Semantic Textual Similarity (STS) and Paraphrase Detection.

### 3.3.2 Architecture and Embedding Strategy

Each sentence is independently passed through a shared BERT encoder. Token embeddings from the final hidden layer are then aggregated using **MEAN pooling** to form fixed-length sentence vectors. Padding tokens are excluded using an attention mask. The resulting sentence embeddings are denoted by  $u$  and  $v$ .

### 3.3.3 Paraphrase Detection (Classification Task)

**Objective:** Given a pair of sentences, determine whether they are paraphrases. This is a binary classification problem.

**Embedding Interaction:** Sentence embeddings  $u$  and  $v$  are combined into a joint representation using:

$$[u, v, |u - v|]$$

This captures both content and contrastive features.

**Prediction Function:** In the original SBERT paper, the prediction function is defined using softmax:

$$o = \text{softmax}(W^T[u, v, |u - v|])$$

This works well for multi-class tasks such as SNLI, where the outputs are entailment, contradiction, and neutral. However, in our case — paraphrase detection is a **binary classification problem**. Therefore, we use the sigmoid function instead of softmax, resulting in:

$$o = \sigma(W^T[u, v, |u - v|])$$

Where  $\sigma$  is the sigmoid activation, and  $W$  is a learned weight matrix.

**Why Sigmoid Instead of Softmax:** Softmax normalizes across multiple output classes and is better suited for multi-class classification. Sigmoid is appropriate for binary classification because it outputs a single scalar probability, directly indicating the likelihood that the sentence pair is a paraphrase. It aligns naturally with the Binary Cross Entropy loss, making it both a computational and conceptual fit.

**Loss Function:** We use Binary Cross Entropy loss for optimization:

$$\mathcal{L}_{\text{classification}} = -\frac{1}{N} \sum_{i=1}^N [y_i \log(o_i) + (1 - y_i) \log(1 - o_i)]$$

Where  $y_i$  is the true label and  $o_i$  is the predicted probability for the  $i$ -th pair.

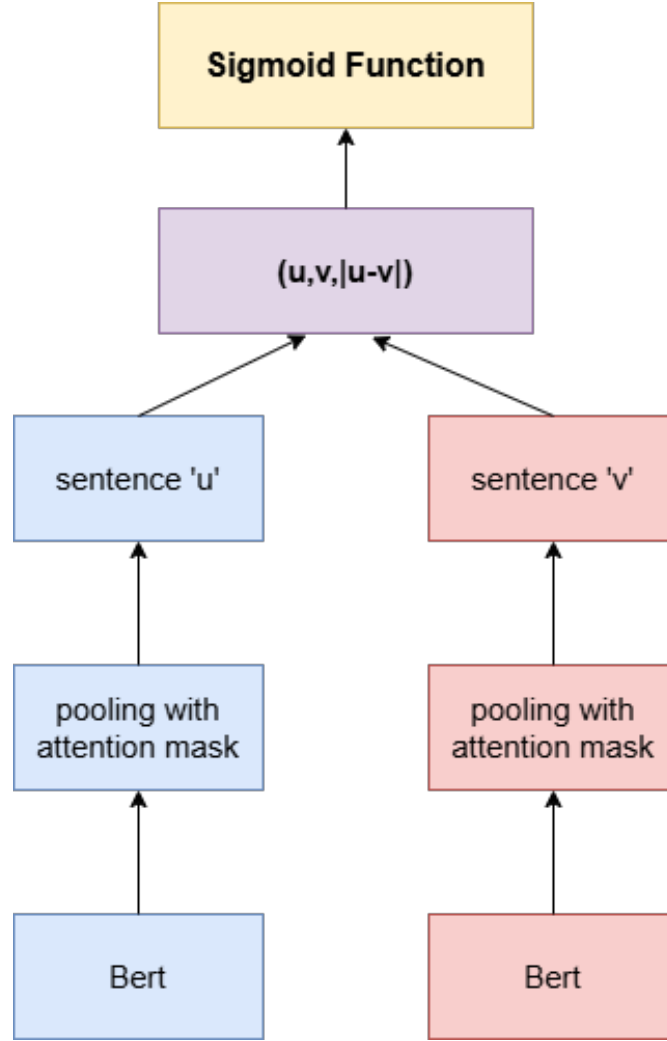


FIGURE 3.1: SBERT Implementation Architecture for Paraphrase Detection. Referenced from *Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks* by Reimers and Gurevych [9]

### 3.3.4 Semantic Textual Similarity (Regression Task)

**Objective:** Estimate a continuous similarity score (0 to 5) that reflects how semantically close two sentences are.

**Embedding Interaction:** Sentence embeddings  $u$  and  $v$  are obtained using MEAN pooling and attention masking, as in the classification task.

**Similarity Computation:** Cosine similarity is computed between the two vectors:

$$\text{sim}(u, v) = \frac{u^\top v}{\|u\| \cdot \|v\|}$$

To match the STSb dataset scale (0 to 5), we use:

$$\text{similarity} = \text{sim}(u, v) \times 2.5 + 2.5$$

**Loss Function:** Mean Squared Error (MSE) loss is employed during training for this task:

$$\mathcal{L}_{\text{regression}} = \frac{1}{N} \sum_{i=1}^N (s_i - \hat{s}_i)^2$$

Where  $s_i$  is the ground truth similarity score and  $\hat{s}_i$  is the predicted similarity.

**Evaluation Metric - Pearson Correlation:** Pearson Correlation measures the linear alignment between predicted and ground truth scores. Formula is

$$r = \frac{\sum (x_i - \bar{x})(y_i - \bar{y})}{\sqrt{\sum (x_i - \bar{x})^2} \cdot \sqrt{\sum (y_i - \bar{y})^2}}$$

where  $x_i$  denotes the predicted score,  $\bar{x}$  is the mean of the predicted scores,  $y_i$  is the ground truth score,  $\bar{y}$  is the mean of the ground truth scores, and  $r$  is the Pearson correlation coefficient.

Unlike MSE, Pearson correlation evaluates how well the model preserves the ranking of similarity across examples — which is the actual goal in STS.

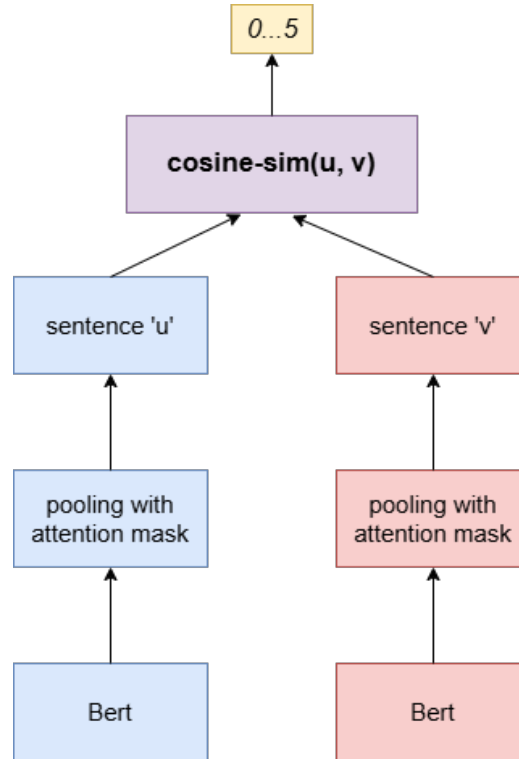


FIGURE 3.2: SBERT Implementation Architecture for Semantic Textual Similarity. Referenced from *Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks* by Reimers and Gurevych [9]

### 3.3.5 Training Details

The paraphrase detection model is trained using the QQP dataset with sigmoid activation and Binary Cross Entropy loss. The STS model is trained using the STSb dataset with cosine similarity and MSE loss. Dropout regularization is applied after pooling to prevent overfitting. Attention masking is consistently used to exclude padding tokens during pooling. Fine-tuning parameters such as learning rate and batch size are selected empirically to optimize performance.

## Chapter 4

# Experimentation (Rooshan Khan)

### 4.1 Phase 1

Phase 1 of our project focused on fine-tuning BERT for the sentiment analysis task using the SST-5 and CFIMDB datasets. This phase aimed to understand the process of fine-tuning for a single task. We experimented with two fine-tuning strategies: (i) fine-tuning only the last linear layer, and (ii) fine-tuning the entire model. The development accuracies for both datasets under each strategy are presented in Table 4.1.

TABLE 4.1: Quantitative Result of Phase 1. Accuracy comparison of BERT on SST and CFIMDB.

	Dev Accuracies			
	Ours		Reference	
	SST	CFIMDB	SST	CFIMDB
last linear layer finetuning	0.409	0.788	0.390	0.780
full model finetuning	0.524	0.967	0.515	0.966

### 4.2 Phase 2

Imagine that we need separate models for different tasks. The biggest issue in this case is memory consumption. Storing 110 million parameters for each model becomes highly inefficient— $n$  tasks would require  $n$  separate models, resulting in a total of  $110n$  million parameters.

Why not use the same BERT model for all tasks? The question is: how? One approach is to share the 12 BERT layers across all tasks while using separate task-specific heads. This idea is illustrated in Figure 4.1.



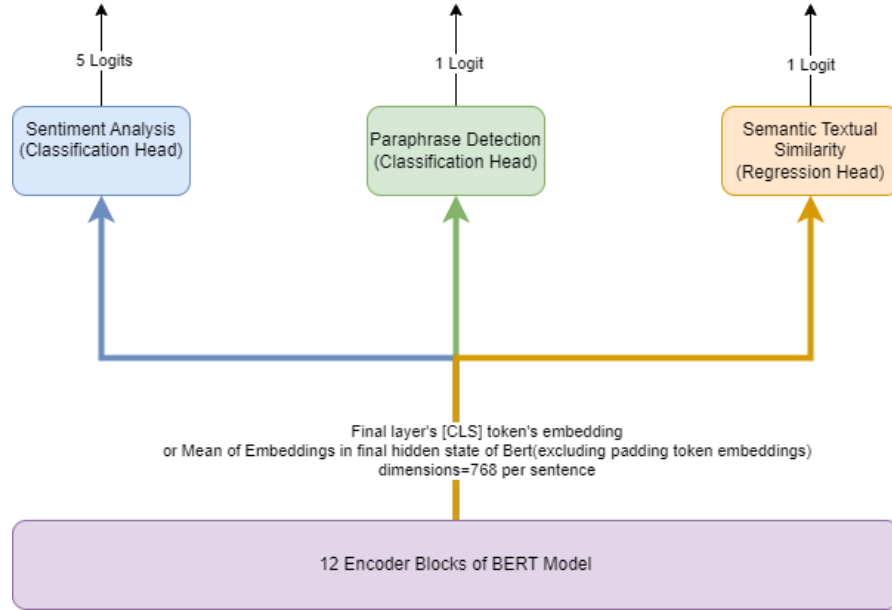


FIGURE 4.1: Proposed Model

We used Colab's L4 GPU for training and evaluating our models. Initially, we adopted a `batch_size=32`, which worked well for baseline training. However, the introduction of SMART regularization significantly increased memory usage, leading to CUDA out-of-memory (OOM) errors.

For the Sentiment Analysis task, we were able to apply SMART successfully with `batch_size=32` without encountering memory issues. However, for the Semantic Textual Similarity (STS) task, training with SMART triggered CUDA OOM errors, prompting us to lower the batch size to 16.

Later, during training for the Paraphrase Detection task, we found that even `batch_size=16` caused CUDA OOM errors. As a result, we further reduced the batch size to 8 for this task to ensure successful training.

The Architecture that we will be using throughout our experimentation is given in Figure 4.2

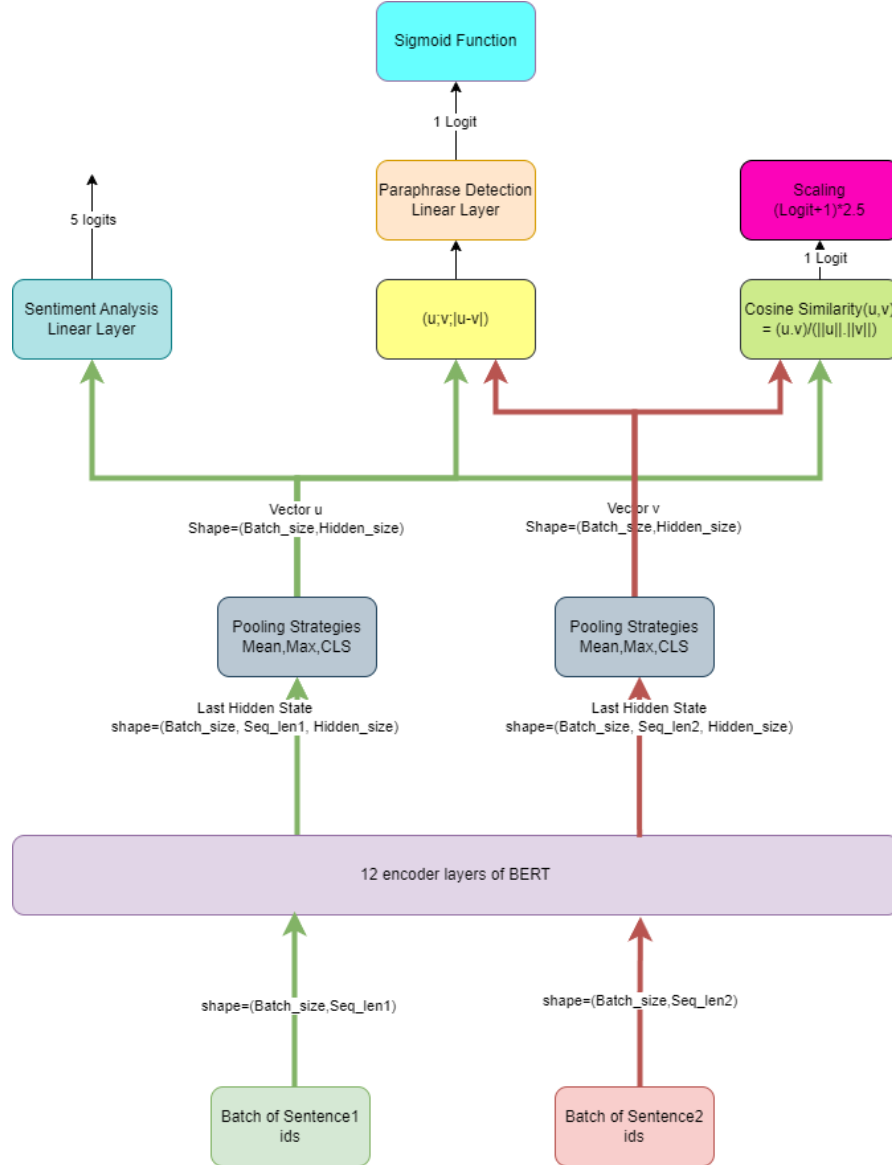


FIGURE 4.2: Final Model Architecture

#### 4.2.1 Individually fine-tuned BERT models

We first trained and tested our models on individual tasks. The best accuracy on the SST dataset was achieved using the mean of all vectors from the last hidden state corresponding to the input words (excluding padding tokens). We employed the SMART algorithm for training. The highest accuracy obtained was **0.533** at **epoch 7**. We set the hyperparameters as follows:  $\lambda_s = 10$  and  $\mu = 1$ . A dropout rate of 0.3 was used during training. Look at Figure 4.3.




FIGURE 4.3: SST train vs dev accuracy over 10 epochs

When we trained our model on the Semantic Textual Similarity (STS) task alone, we achieved a best Pearson correlation of 0.821 on the development set. We achieved Pearson correlation on the first epoch. We used SBERT architecture with mean pooling for regression tasks that finds cosine similarity between sentences. We shifted and scaled this cosine similarity to be in the range from 0 to 5 instead of -1 to 1. We also used the SMART algorithm during training. We used  $\lambda_s = 5$  and  $\mu = 1$ . A dropout rate of 0.3 was used during training. As shown in Figure 4.4, the dev STS correlation starts to drop after epoch 0, indicating potential overfitting.

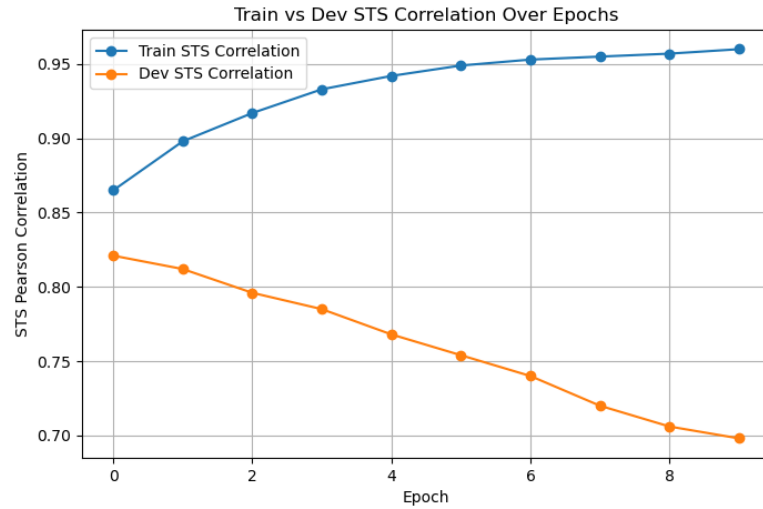


FIGURE 4.4: STS train vs dev correlation over 10 epochs

For the Paraphrase Detection task, our model achieved a maximum accuracy of 0.875 with hyperparameters set to  $\lambda_s = 5$  and  $\mu = 1$ . Due to computational constraints, we

trained the model for only one epoch. Specifically, incorporating SMART significantly increased training time to approximately 5 hours per epoch, which was not feasible given our limited compute resources. A dropout rate of 0.1 was used during training. Batch size was set to 8 to avoid CUDA out-of-memory error. We used SBERT architecture with mean pooling for this classification task.

The results for separate models trained individually for each task are summarized in Table 4.2.

Task	Metric	Value
Sentiment Classification (SST-5)	Accuracy	0.533
Paraphrase Detection (QQP)	Accuracy	0.875
Semantic Textual Similarity (STS-B)	Pearson Correlation	0.821

TABLE 4.2: Best performance of separate models fine-tuned individually for each task

## 4.2.2 Accuracies and Pearson Correlation for Fine-Tuning the same model on Multiple Tasks

### 4.2.2.1 Sequential Fine-Tuning: SBERT Only

We trained the model for **3** epochs on the Paraphrase Detection task, **10** epochs on the Semantic Textual Similarity (STS) task, and **10** epochs on the Sentiment Analysis task.

Task	Metric	Value
Paraphrase Detection	Accuracy	0.6285
Semantic Textual Similarity (STS)	Pearson Correlation	0.5444
Sentiment Analysis (SST)	Accuracy	0.1698

TABLE 4.3: Development set performance before fine-tuning

After training for 3 epochs on the QQP dataset, the best Paraphrase Detection accuracy achieved was **0.7441**. See Figure 4.5 for the training and development curves.




FIGURE 4.5: Train vs Dev Accuracy for Paraphrase Detection using SBERT Only (3 epochs)

Task	Metric	Value
Paraphrase Detection	Accuracy	0.7441
Semantic Textual Similarity (STS)	Pearson Correlation	0.7127
Sentiment Analysis (SST)	Accuracy	0.1735

TABLE 4.4: Development set performance after fine-tuning on Paraphrase Detection for 3 epochs

Subsequently, we fine-tuned the best-performing model for **10** epochs on the STS task.




FIGURE 4.6: Train vs Dev Correlation for Semantic Textual Similarity using SBERT Only (10 epochs)

Task	Metric	Value
Paraphrase Detection	Accuracy	0.8353
Semantic Textual Similarity (STS)	Pearson Correlation	0.6890
Sentiment Analysis (SST)	Accuracy	0.1835

TABLE 4.5: Development set performance after fine-tuning on the STS task

Finally, we fine-tuned the best-performing model for **10** epochs on the Sentiment Analysis task.




FIGURE 4.7: Train vs Dev Accuracy for Sentiment Analysis using SBERT Only (10 epochs)

Task	Metric	Value
Paraphrase Detection	Accuracy	0.7940
Semantic Textual Similarity (STS)	Pearson Correlation	0.6220
Sentiment Analysis (SST)	Accuracy	0.4980

TABLE 4.6: Development set performance after fine-tuning for 10 epochs on SST

- Observations:** We observed that training on each task could negatively impact the performance on dissimilar tasks. For example, training on the SST task hurt the model’s performance on the STS task. In contrast, training on a particular task improved performance on related tasks. For instance, fine-tuning the model for paraphrase detection improved the correlation score on the STS task, while training on the STS task enhanced the accuracy on the paraphrase detection task.

#### 4.2.2.2 Sequential Fine-Tuning: SBERT+SMART

For this training paradigm, we were constrained to use a batch size of 8 due to CUDA out-of-memory errors for higher batch sizes. The SMART regularization technique significantly slowed down training. Specifically, the paraphrase detection task failed for batch sizes above 8. Hence, we decided to use a uniform batch size of 8 across all tasks.

Fine-tuning for paraphrase detection (QQP) took approximately 5 hours per epoch. Due to limited computational resources, we fine-tuned for only one epoch. The development accuracy obtained after this epoch was 0.807. The resulting metrics for all three tasks are summarized in Table 4.7.

Task	Metric	Value
Paraphrase Detection	Accuracy	0.8067
Semantic Textual Similarity (STS)	Pearson Correlation	0.7567
Sentiment Analysis (SST)	Accuracy	0.2089

TABLE 4.7: Development set performance after fine-tuning on Paraphrase Detection for 1 epoch

Subsequently, we fine-tuned the model on the STS dataset for 10 epochs. Interestingly, we observed that the highest Pearson correlation was achieved in the first epoch, followed by a consistent decline. This trend is shown in Figure 4.8.

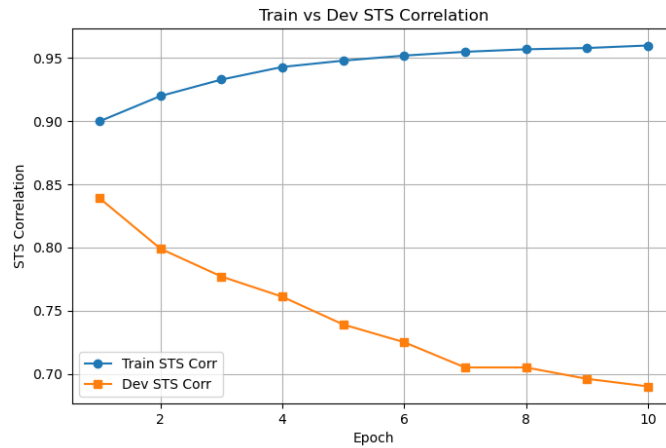


FIGURE 4.8: Train vs Dev Pearson Correlation for STS using SMART SBERT (10 epochs)

The post-training metrics for all tasks are listed below.

Task	Metric	Value
Paraphrase Detection	Accuracy	0.8045
Semantic Textual Similarity (STS)	Pearson Correlation	0.8394
Sentiment Analysis (SST)	Accuracy	0.2044

TABLE 4.8: Development set performance after fine-tuning for 10 epochs on STS

Finally, we fine-tuned the best model from the STS stage on the SST dataset for 10 epochs. The train vs. dev accuracy trends are illustrated in Figure 4.9.



FIGURE 4.9: Train vs Dev Accuracy for SST using SMART SBERT (10 epochs)

The resulting performance metrics are reported in Table 4.9.

Task	Metric	Value
Paraphrase Detection	Accuracy	0.7350
Semantic Textual Similarity (STS)	Pearson Correlation	0.7880
Sentiment Analysis (SST)	Accuracy	0.5210

TABLE 4.9: Development set performance after fine-tuning for 10 epochs on SST

We changed the dropout to 0.1 and found that after training for 1 epoch on the paraphrase detection task, the accuracy reached 0.8749. However, when we trained this model on the STS task, the accuracy for STS rose to 0.8580, but the paraphrase detection accuracy dropped to 0.7523. This was due to catastrophic forgetting. Therefore, we decided to train the model as follows:

1. Train the model for one epoch on the paraphrase task using the complete dataset.
2. Train the model for one epoch on the STS dataset (full dataset).
3. Train the model for one epoch on the SST dataset (full dataset).



4. Train the model on the paraphrase task using only 3% of the dataset.

We repeat steps 2 to 4 ten times and save the best model. Overall performance is calculated after training on each task using the following formula:

$$\text{Overall performance} = \frac{\left(\frac{\text{dev\_sts\_corr}+1}{2} + \text{sst\_dev\_acc} + \text{paraphrase\_dev\_acc}\right)}{3}$$

This updated training procedure can be seen in Figure 4.10.

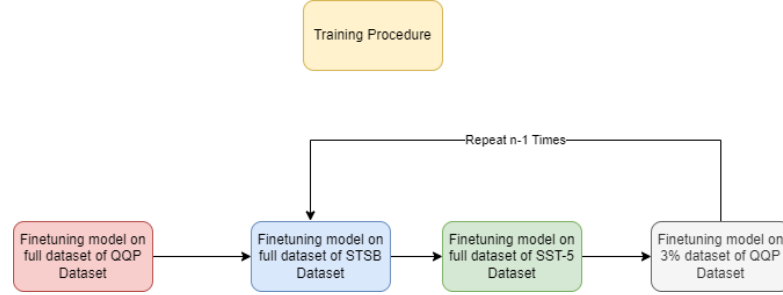


FIGURE 4.10: Updated Training Procedure

The optimal model is the one with the highest overall performance across all tasks.

The following table shows the metrics for our best model so far:

Task	Metric	Score
Sentiment Classification	Accuracy	0.537
	F1 Score	0.528
Paraphrase Detection	Accuracy	0.864
	F1 Score	0.864
Semantic Textual Similarity	Pearson Correlation	0.819

TABLE 4.10: Evaluation metrics for each task (SMART+SBERT)

Since we changed the training procedure and dropout (from 0.3 to 0.1), it seems logical to retrain the **SBERT\_Only** model using the new procedure and dropout for a fair comparison.

We obtain the following results:

Task	Metric	Score
Sentiment Classification	Accuracy	0.501
	F1 Score	0.495
Paraphrase Detection	Accuracy	0.854
	F1 Score	0.856
Semantic Textual Similarity	Pearson Correlation	0.680

TABLE 4.11: Evaluation metrics for each task (SBERT Only)

#### 4.2.2.3 Sequential Fine-Tuning: SBERT+SimCSE+SMART

Before applying the training procedure shown in Figure 4.10, we performed contrastive loss minimization (Supervised SimCSE Loss) on the SNLI dataset for 3 epochs, as recommended in *SimCSE: Simple Contrastive Learning of Sentence Embeddings* by Gao et al. [2]. The best model was selected based on the highest Pearson correlation on the STS-B task. Sentence embeddings were computed as the mean of the token embeddings from the last layer of BERT. We obtain the following results after this step:

Task	Metric	Score
Sentiment Classification	Accuracy	0.221
	F1 Score	0.174
Paraphrase Detection	Accuracy	0.626
	F1 Score	0.503
Semantic Textual Similarity	Pearson Correlation	0.812

TABLE 4.12: Evaluation metrics for each task

Then after this we begin the training procedure of Figure 4.10. When we train the model for 1 complete QQP dataset we get an accuracy of 0.876. However, with the later training we could reach a maximum performance of 0.764. This is lower than what we got for SMART SBERT (Our best model so far). The results are shown in the table below.

Task	Metric	Score
Sentiment Classification	Accuracy	0.507
	F1 Score	0.494
Paraphrase Detection	Accuracy	0.864
	F1 Score	0.863
Semantic Textual Similarity	Pearson Correlation	0.843
<b>Overall Performance</b>		<b>0.764</b>

TABLE 4.13: Evaluation metrics for each task

### 4.3 Analysis of the Results

Our best-performing model achieved an overall score of 0.770%. These results are in line with the top-performing individual models (developed by us) discussed in Section 4.2.1. Additionally, our outcomes are competitive with state-of-the-art results reported for each task. The benchmark scores listed in Table 4.14 are sourced from *Papers with Code* [6].

Task	Metric	Score
Sentiment Classification	Accuracy	0.532
Paraphrase Detection	Accuracy	0.910
Semantic Textual Similarity	Pearson Correlation	0.8479
<b>Overall Performance</b>		0.7887

TABLE 4.14: State-of-the-art BERT-base results for each task

We have two strong models: SMART SBERT (SS) and SMART SBERT with SimCSE (SSS). Both SS and SSS perform approximately equally well on the paraphrase detection task. However, SS performs better on the SST-5 task, while SSS performs better on the STS-B task. The superior performance of SSS on STS-B is understandable, as the supervised SimCSE loss is specifically optimized to enhance performance on semantic textual similarity tasks such as STS-B. Depending on the task priority, either of the two models can be selected.

## Chapter 5

# Possible Improvements (Areesha Noor)

The Quora Question Pairs (QQP) dataset used is imbalanced, containing approximately 63% of the examples labeled as non-duplicates and 37% as duplicates, in both training and validation splits. Despite this imbalance, we used **accuracy** as the primary evaluation metric and based our model selection on it. This is mainly to make a fair comparison among our models, as state-of-the-art results also use this metric.

To better account for class imbalance, the accuracy metric may not be a reliable metric as it provides an overall performance score. Instead, we can focus on computing the **sample-weighted F1 score**.

Similar to the Quora Question Pairs dataset, the SST-5 dataset is also imbalanced, where the class distribution changes significantly across five sentiment categories. In the training split dataset, the most frequent classes are *Positive* and *Negative*, making up 27.4% and 26% of the total dataset, respectively. Meanwhile, *Very Positive* and *Very Negative* classes contain only 15.1% and 12.8% of the total data, respectively.

Given the imbalanced class distribution for the above datasets, using the accuracy metric may not provide an accurate assessment of a model's performance, as it could heavily rely on majority classes. So, like the QQP dataset, here also, the **sample-weighted F1 score** could be used for the SST-5 dataset to better evaluate performance across all classes. It considers two competing metrics: precision and recall.

$$\text{Precision} = \frac{TP}{TP + FP}$$
$$\text{Recall} = \frac{TP}{TP + FN}$$

where:

- TP is the number of true positives
- FP is the number of false positives
- FN is the number of false negatives

Sample F1 score helps by weighing the F1 score for each class on the basis of its true instances. First, we compute class-wise F1 scores as:

$$F1 = \frac{2 \times \text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}}$$

Then, we take the weighted average of F1 scores for all classes as:

$$\text{Weighted F1 Score} = \sum_{i=1}^N w_i \times F1_i$$

where:

$$w_i = \frac{\text{Number of samples in class } i}{\text{Total number of samples}}$$

Here, the sample-weighted F1 score provides a more balanced evaluation by equitably weighing each class and thus proves more suitable for imbalanced datasets like SST-5.

In our Python implementation, we modified one parameter in the evaluation function by changing:

```
f1_score(y_true, y_pred, average='macro')
```

to

```
f1_score(y_true, y_pred, average='weighted')
```

This computes class-wise F1 scores and then multiplies them by their weights (number of true instances). This approach helps in providing more realistic performance evaluation in the case of imbalanced datasets like SST-5 or Quora Question Pairs.

## Appendix A

# Implementation of Techniques in Python (Rooshan Khan)

The foundational codebase utilized in our final year project was adapted from materials provided in the Stanford *CS224n: Natural Language Processing with Deep Learning* course [4, 10].

### A.1 SMART

The theory related to SMART was explained in Section 3.1. Now let's talk about how we implemented the SMART algorithm in python.

The Listing A.1 shows the code that trains model on the SST dataset using SMART. It correctly implements the SMART algorithm described in the research paper *SMART: Robust and Efficient Fine-Tuning for Pre-trained Natural Language Models through Principled Regularized Optimization* by Jiang et al. [3]. Hyperparameters are set manually. An `if` condition for  $\lambda_s = 0$  or  $\mu = 0$  has been added to easily switch between training with SMART and without SMART. We keep  $S = 1$  and  $T_x = 1$  throughout our training on all tasks. This choice is motivated by computational complexity. Moreover, the authors of the paper themselves used and recommended setting  $S = 1$  and  $T_x = 1$ . You can see that the `predict_sentiment` method is called three times in the loop for every batch. The reason is as follows: it is first needed inside the  $T_x$ -loop to compute  $\tilde{g}_i$ . It is then needed a second time outside the  $T_x$ -loop but inside the  $S$ -loop, where it is used to compute the logits without noise and the logits with updated noise; these logits are used to calculate  $R_s$ . Finally, the `predict_sentiment` method is called a third time to compute  $D_{\text{Breg}}$ . In the third case, we use `model_copy.predict_sentiment` because `model_copy` has parameters  $\tilde{\theta}_t$ .

```
# Values of Few hyperparameters

epsilon=1e-5
eps = 1e-8
```



```

        max_ls, _ = torch.max(ls, -1)
        ls = ls.sum()
        Rs = torch.sum(max_ls) / max_ls.shape[0]
        Embedding_in_wn
        ls.backward()
        gi = Embedding_in_wn.grad / Embedding_in_wn.shape
[0]
        gi_tilde = gi / (torch.norm(gi, p=float('inf'), dim
=(1, 2)).view(gi.shape[0], 1, 1) + eps)
        noise = noise.clone() + eta * gi_tilde.clone()
        logits, logits_wn, Embedding_in, Embedding_in_wn =
model.predict_sentiment(b_ids, b_mask, noise)
        P1 = F.softmax(logits, dim=1)
        Q = F.softmax(logits_wn, dim=1)
        KL_divergence_PQ = P1 * torch.log((P1 + eps) / (Q
+ eps))
        KL_divergence_QP = Q * torch.log((Q + eps) / (P1 +
eps))
        ls = KL_divergence_PQ + KL_divergence_QP
        max_ls, _ = torch.max(ls, -1)
        ls = ls.sum()
        Rs = torch.sum(max_ls) / max_ls.shape[0]
        logits2, logits2_wn, Embedding_in_2,
Embedding_in_wn_2 = model_copy.predict_sentiment(b_ids, b_mask,
noise)
        P2 = F.softmax(logits2, dim=1)
        KL_divergence_P1P2 = P1 * torch.log((P1 + eps) / (
P2 + eps))
        KL_divergence_P2P1 = P2 * torch.log((P2 + eps) / (
P1 + eps))
        ls2 = KL_divergence_P1P2 + KL_divergence_P2P1
        D_Breg = ls2.sum() / args.batch_size
        optimizer.zero_grad()
        loss = F.cross_entropy(logits, b_labels.view(-1),
reduction='sum') / args.batch_size + lambda_s * Rs + mu * D_Breg
        loss.backward()
        optimizer.step()
        theta_tilde_s = copy.deepcopy(model.state_dict())
        theta_t = copy.deepcopy(theta_tilde_s)
        theta_t_tilde = copy.deepcopy(theta_t)
        for key in theta_tilde_s.keys():
            theta_t_tilde[key] = (1 - Beta) * theta_tilde_s[
key] + Beta * theta_t_tilde[key]
        else:
            optimizer.zero_grad()
            logits = model.predict_sentiment(b_ids, b_mask)

```



```

        loss = F.cross_entropy(logits, b_labels.view(-1),
reduction='sum') / args.batch_size
        loss.backward()
        optimizer.step()
        train_loss += loss.item()
        num_batches += 1

train_loss = train_loss / (num_batches)

sst_train_acc, sst_train_f1, *_ = model_eval_sst(
sst_train_dataloader, model, device)
sst_dev_acc, sst_dev_f1, *_ = model_eval_sst(
sst_dev_dataloader, model, device)

if sst_dev_acc > best_sst_dev_acc:
    best_sst_dev_acc = sst_dev_acc
    save_model(model, optimizer, args, config, args.filepath)

print(f"Epoch {epoch}: train loss :: {train_loss :.3f}, sst
train acc :: {sst_train_acc :.3f}, sst dev acc :: {sst_dev_acc
:.3f}")

```

LISTING A.1: Sentiment Analysis Training with SMART

Listing A.2 shows the code that trains the model on the STS dataset using SMART. This code is similar to that given in Listing A.1, which trains the model for sentiment analysis. There are two differences. The first difference lies in the computation of  $l_s$ . In this case,  $l_s(p, q) = (p - q)^2$ , rather than the sum of the KL divergences, i.e.,  $\text{KL}(P\|Q)$  and  $\text{KL}(Q\|P)$ . The second difference is that, since there are two inputs (embeddings of all tokens in the input sentence), we need to find and update the noise for both inputs.

```

# Values of Few hyperparameters
epsilon=1e-5
eps = 1e-8
lambda_s=5
std=1e-5
mean=0
S=1
Tx=1
eta=1e-3
mu=1

# Run for the specified number of epochs.
for epoch in range(args.epochs//10*2):
    model.train()
    train_loss = 0
    num_batches = 0
    theta_t=copy.deepcopy(model.state_dict())

```

```

theta_tilde_t=copy.deepcopy(theta_t)
Beta=0.99
n=0
for batch in tqdm(sts_train_dataloader, desc=f'train-{
epoch}', disable=TQDM_DISABLE):
    b_ids1, b_mask1,b_ids2, b_mask2, b_labels = (batch['
token_ids_1'], batch['attention_mask_1'],batch['token_ids_2'],
batch['attention_mask_2'],batch['labels'])

    b_ids1 = b_ids1.to(device)
    b_mask1 = b_mask1.to(device)
    b_ids2 = b_ids2.to(device)
    b_mask2 = b_mask2.to(device)
    b_labels=b_labels.to(device)
    if lambda_s!=0 or mu!=0:
        theta_tilde_s=copy.deepcopy(theta_t)

        if n==38:
            Beta=0.999
            n+=1
            model_copy = copy.deepcopy(model)
            model_copy.load_state_dict(theta_tilde_t)
            for s in range(0,S):
                noise1 = torch.randn(b_ids1.shape[0], b_ids1.
shape[1], 768)*std+mean
                noise1=noise1.to(device)
                noise2 = torch.randn(b_ids2.shape[0], b_ids2.
shape[1], 768)*std+mean
                noise2=noise2.to(device)
                for m in range(0,Tx):
                    logit0,logit0_wn,Embedding_in1,
Embedding_in2,Embedding_in1_wn,Embedding_in2_wn = model.
predict_similarity(b_ids1, b_mask1, b_ids2, b_mask2, noise1,
noise2)

                    ls=(logit0_wn-logit0)**2
                    max_ls=ls
                    ls=ls.sum()
                    Rs=torch.sum(max_ls)/max_ls.shape[0]
                    Embedding_in1_wn.retain_grad()
                    Embedding_in2_wn.retain_grad()
                    ls.backward()
                    gi1=Embedding_in1_wn.grad/Embedding_in1_wn
.shape[0]

                    gi2=Embedding_in2_wn.grad/Embedding_in2_wn
.shape[0]

                    gi1_tilde=gi1/(torch.norm(gi1,p=float('inf
'),dim=(1, 2)).view(gi1.shape[0],1,1)+eps)

```

```

        gi2_tilde=gi2/(torch.norm(gi2,p=float('inf'),dim=(1, 2)).view(gi2.shape[0],1,1)+eps)
        noise1 = noise1.clone() + eta * gi1_tilde.clone()
        noise2 = noise2.clone() + eta * gi2_tilde.clone()

        logit,logit_wn,Embedding_in1_1,Embedding_in2_1,Embedding_in1_wn_1,Embedding_in2_wn_1 = model.predict_similarity(b_ids1, b_mask1, b_ids2, b_mask2, noise1, noise2)

        ls=(logit-logit_wn)**2
        max_ls=ls
        ls=ls.sum()
        Rs=torch.sum(max_ls)/max_ls.shape[0]
        logit2,logit2_wn,Embedding_in1_2,Embedding_in2_2,Embedding_in1_wn_2,Embedding_in2_wn_2 = model_copy.predict_similarity(b_ids1, b_mask1, b_ids2, b_mask2, noise1, noise2)

        ls2=(logit2-logit)**2
        D_Breg=ls2.sum()/args.batch_size
        optimizer.zero_grad()
        loss=F.mse_loss(logit, b_labels.view(-1).float(), reduction='sum') / args.batch_size + lambda_s*Rs +mu*D_Breg
        loss.backward()
        optimizer.step()
        theta_tilde_s=copy.deepcopy(model.state_dict())
    )

    theta_t=copy.deepcopy(theta_tilde_s)
    theta_t_tilde = copy.deepcopy(theta_t)
    for key in theta_tilde_s.keys():
        theta_t_tilde[key] = (1 - Beta) * theta_tilde_s[key] + Beta * theta_t_tilde[key]
    else:
        optimizer.zero_grad()
        logit = model.predict_similarity(b_ids1, b_mask1, b_ids2, b_mask2)
        loss = F.mse_loss(logit, b_labels.view(-1).float(), reduction='sum') / args.batch_size
        loss.backward()
        optimizer.step()
        train_loss += loss.item()
        num_batches += 1

    train_loss = train_loss / (num_batches)

    train_sts_corr= model_eval_sts(sts_train_dataloader, model, device)

```

```

        dev_sts_corr= model_eval_sts(sts_dev_dataloader, model,
device)

        if dev_sts_corr > best_sts_corr:
            best_sts_corr = dev_sts_corr
            save_model(model, optimizer, args, config, args.
filepath)

        print(f"Epoch {epoch}: train loss :: {train_loss :.3f},
train_sts_corr :: {train_sts_corr :.3f}, dev_sts_corr :: {
dev_sts_corr :.3f}")

```

LISTING A.2: Sematic Textual Similarity Training with SMART

Listing A.3 shows the code that trains the model on the Quora Question Pairs (QQP) dataset using SMART. This code is similar to that given in Listing A.1. Although QQP is also a classification task like SST, it involves binary classification. There is a single logit output. Thus, we cannot apply the softmax function to the logit. Instead, we first compute  $\text{sigmoid}(\text{logit})$  and  $\text{sigmoid}(\text{logit\_wn})$ , treating these values as the probabilities of being a paraphrase. The probabilities of not being a paraphrase are computed as  $1 - \text{sigmoid}(\text{logit})$  and  $1 - \text{sigmoid}(\text{logit\_wn})$ .

```

# Values of Few hyperparameters
epsilon=1e-5
eps = 1e-8
lambda_s=5
std=1e-5
mean=0
S=1
Tx=1
eta=1e-3
mu=1

# Run for the specified number of epochs.
for epoch in range((args.epochs//10)*3):
    model.train()
    train_loss = 0
    num_batches = 0
    theta_t=copy.deepcopy(model.state_dict())
    theta_tilde_t=copy.deepcopy(theta_t)
    Beta=0.99
    n=0
    for batch in tqdm(para_train_dataloader, desc=f'train-{
epoch}', disable=TQDM_DISABLE):
        b_ids1, b_mask1,b_ids2, b_mask2,b_labels= (batch['
token_ids_1'], batch['attention_mask_1'],batch['token_ids_2'],
batch['attention_mask_2'],batch['labels'])

```

```

        b_ids1 = b_ids1.to(device)
        b_mask1 = b_mask1.to(device)
        b_ids2 = b_ids2.to(device)
        b_mask2 = b_mask2.to(device)
        b_labels=b_labels.to(device)

    if lambda_s!=0 or mu!=0:
        theta_tilde_s=copy.deepcopy(theta_t)
        if n==1769:
            Beta=0.999
            n+=1
            model_copy = copy.deepcopy(model)
            model_copy.load_state_dict(theta_tilde_t)
            for s in range(0,S):
                noise1 = torch.randn(b_ids1.shape[0], b_ids1.
shape[1], 768)*std+mean
                noise1=noise1.to(device)
                noise2 = torch.randn(b_ids2.shape[0], b_ids2.
shape[1], 768)*std+mean
                noise2=noise2.to(device)
                for m in range(0,Tx):
                    logit,logit_wn,_,_,Embedding_in1_wn,
Embedding_in2_wn = model.predict_paraphrase(b_ids1, b_mask1,
b_ids2, b_mask2, noise1, noise2)
                    prob = torch.sigmoid(logit)      # shape: (
batch_size, 1)
                    prob_wn = torch.sigmoid(logit_wn) # with
noise
                    P1 = torch.cat([1 - prob, prob], dim=1) #
[P(not paraphrase), P(paraphrase)]
                    Q = torch.cat([1 - prob_wn, prob_wn], dim
=1)

                    KL_divergence_PQ = P1 * torch.log((P1 +
eps) / (Q + eps))
                    KL_divergence_QP = Q * torch.log((Q + eps)
/ (P1 + eps))

                    ls=KL_divergence_PQ+KL_divergence_QP
                    ls=ls.sum()
                    ls.backward()
                    gi1_tilde=Embedding_in1_wn.grad/
Embedding_in1_wn.shape[0]
                    gi2_tilde=Embedding_in2_wn.grad/
Embedding_in2_wn.shape[0]
                    gi1_tilde/=(torch.norm(gi1_tilde,p=float('
inf'),dim=(1, 2)).view(gi1_tilde.shape[0],1,1)+eps)

```

```

        gi2_tilde/=(torch.norm(gi2_tilde,p=float('
inf'),dim=(1, 2)).view(gi2_tilde.shape[0],1,1)+eps)
        noise1 += eta * gi1_tilde
        noise2 += eta * gi2_tilde
        logit,logit_wn,*rest = model.
predict_paraphrase(b_ids1, b_mask1, b_ids2, b_mask2, noise1,
noise2)

        prob = torch.sigmoid(logit)      # shape: (
batch_size, 1)
        prob_wn = torch.sigmoid(logit_wn) # with
noise

        P1 = torch.cat([1 - prob, prob], dim=1) # [P(
not paraphrase), P(paraphrase)]
        Q = torch.cat([1 - prob_wn, prob_wn], dim=1)
        KL_divergence_PQ = P1 * torch.log((P1 + eps) /
(Q + eps))

        KL_divergence_QP = Q * torch.log((Q + eps) / (
P1 + eps))

        ls=KL_divergence_PQ+KL_divergence_QP
        max_ls,_=torch.max(ls,-1)
        ls=ls.sum()
        Rs=torch.sum(max_ls)/max_ls.shape[0]
        logit2,*rest = model_copy.predict_paraphrase(
b_ids1, b_mask1, b_ids2, b_mask2, noise1, noise2)
        prob2 = torch.sigmoid(logit2)      # shape: (
batch_size, 1)

        P2 = torch.cat([1 - prob2, prob2], dim=1) # [
P(not paraphrase), P(paraphrase)]
        KL_divergence_P1P2 = P1 * torch.log((P1 + eps)
/ (P2 + eps))

        KL_divergence_P2P1 = P2 * torch.log((P2 + eps)
/ (P1 + eps))

        ls2=KL_divergence_P1P2+KL_divergence_P2P1
        D_Breg=ls2.sum()/args.batch_size
        optimizer.zero_grad()
        loss=model.paraphrase_simcse_loss(b_ids1,
b_mask1,b_ids2, b_mask2,b_labels,pooling='mean',temperature
=0.05) + lambda_s*Rs + mu*D_Breg
        loss.backward()
        optimizer.step()
        theta_tilde_s=copy.deepcopy(model.state_dict()
)

        theta_t=copy.deepcopy(theta_tilde_s)
        theta_t_tilde = copy.deepcopy(theta_t)
        for key in theta_tilde_s.keys():
            theta_t_tilde[key] = (1 - Beta) *
theta_tilde_s[key] + Beta * theta_t_tilde[key]

```

```

        else:
            optimizer.zero_grad()
            logit = model.predict_similarity(b_ids1, b_mask1,
            b_ids2, b_mask2)
            loss = F.binary_cross_entropy_with_logits(logit.
            squeeze(), b_labels.view(-1).float(), reduction='sum') / args.
            batch_size
            loss.backward()
            optimizer.step()
            train_loss += loss.item()
            num_batches += 1
        train_loss = train_loss / (num_batches)

        paraphrase_train_acc = model_eval_paraphrase(
        para_train_dataloader, model, device)
        paraphrase_dev_acc = model_eval_paraphrase(
        para_dev_dataloader, model, device)

        if paraphrase_dev_acc > best_paraphrase_dev_acc:
            best_paraphrase_dev_acc = paraphrase_dev_acc
            save_model(model, optimizer, args, config, args.
            filepath)

        print(f"Epoch {epoch}: train loss :: {train_loss :.3f},
        paraphrase train acc :: {paraphrase_train_acc :.3f}, paraphrase
        dev acc :: {paraphrase_dev_acc :.3f}")

```

LISTING A.3: Paraphrase Detection Training with SMART

### A.1.1 Implementation of methods for predictions

Look at the codes shown in Listing A.5, Listing A.6, and Listing A.7. All of these use mean pooling as the default. The `if` condition checking whether `noise` is a tensor or not was added for the implementation of SMART during training. The default value of `noise` is 0, which is an integer and not a tensor. This allows the same `predict` method to be called during both training and evaluation. During evaluation, `noise` remains an integer, so SMART is not applied. To achieve this behavior, we also modified the `forward` method of `BERTModel`, as shown in Listing A.4. The dictionary returned by this method differs between training with SMART and evaluation. Another important detail in Listings A.5, A.6, and A.7 is that mean pooling is not performed using `torch.mean(tensor, dim=-1)`. Upon inspecting `b_ids`, we observed the presence of padding tokens. If we had used simple mean pooling with `torch.mean`, it would have incorrectly averaged the embeddings of all tokens, including the padding tokens, leading to distorted representations. Instead, we applied mean pooling using the attention mask to sum only the embeddings corresponding to the [CLS] token, [SEP] token, and the

actual word tokens between them, while ignoring the padding tokens. The same masked mean pooling approach is consistently used across all the `predict` methods.

```
def forward(self, input_ids, attention_mask, perturbation=0,
epsilon=1e-5):
    """
    input_ids: [batch_size, seq_len], seq_len is the max
length of the batch
    attention_mask: same size as input_ids, 1 represents non-
padding tokens, 0 represents padding tokens
    """
    # Get the embedding for each input token.
    embedding_output = self.embed(input_ids=input_ids).
requires_grad_()
    # Feed to a transformer (a stack of BertLayers).
    sequence_output = self.encode(embedding_output,
attention_mask=attention_mask)

    # Get cls token hidden state.
    first_tk = sequence_output[:, 0]
    first_tk = self.pooler_dense(first_tk)
    first_tk = self.pooler_af(first_tk)
    if isinstance(perturbation, torch.Tensor):
        norm = (torch.norm(perturbation, p=float('inf'), dim
=(1, 2))).view(perturbation.shape[0], 1, 1)
        perturbation=perturbation*epsilon/(norm+1e-8)
        embedding_output_with_noise=embedding_output+
perturbation
        embedding_output_with_noise=
embedding_output_with_noise.detach().requires_grad_()
        sequence_output_n = self.encode(
embedding_output_with_noise, attention_mask=attention_mask)
        # Get cls token hidden state for
embedding_output_with_noise.
        first_tk_n = sequence_output_n[:, 0]
        first_tk_n = self.pooler_dense(first_tk_n)
        first_tk_n = self.pooler_af(first_tk_n)
        return {'last_hidden_state': sequence_output, '
pooler_output': first_tk, 'embedding_output': embedding_output, '
last_hidden_state_wn': sequence_output_n, 'pooler_output_wn':
first_tk_n, 'embedding_output_with_noise':
embedding_output_with_noise}
    return {'last_hidden_state': sequence_output, '
pooler_output': first_tk}
```

LISTING A.4: Changed forward method in BERTModel



```

def predict_sentiment(self, input_ids, attention_mask, noise=0,
pooling='mean'):
    '''Given a batch of sentences, outputs logits for classifying
    sentiment.
    There are 5 sentiment classes:
    (0 - negative, 1- somewhat negative, 2- neutral, 3- somewhat
    positive, 4- positive)
    Thus, your output should contain 5 logits for each sentence.
    '''
    ### TODO
    embedding_output=self.forward(input_ids, attention_mask, noise=
noise)
    if pooling=='mean':
        hidden_states=embedding_output['last_hidden_state']
        mask=attention_mask.unsqueeze(-1).expand(-1, -1,
hidden_states.shape[2]).float()
        den=torch.sum(attention_mask, dim=1).unsqueeze(-1).expand
(-1, hidden_states.shape[2])+1e-9
        vector = torch.sum(mask*(embedding_output['
last_hidden_state']), dim=1)/den
    elif pooling=='cls':
        vector=embedding_output['pooler_output']
        Dropout_layer_output=self.dropout_layer(vector)
        logits=self.sentiment_classifier_layer(Dropout_layer_output)
    if isinstance(noise, torch.Tensor):
        if pooling=='mean':
            vector_wn = torch.sum(mask*(embedding_output['
last_hidden_state_wn']), dim=1)/den
        elif pooling=='cls':
            vector_wn=embedding_output['pooler_output_wn']
            Dropout_layer_output_wn=self.dropout_layer(vector_wn)
            logits_wn=self.sentiment_classifier_layer(
Dropout_layer_output_wn)
            Embedding_in=embedding_output['embedding_output']
            Embedding_in_wn=embedding_output['
embedding_output_with_noise']

        return logits, logits_wn, Embedding_in, Embedding_in_wn
    return logits

```

LISTING A.5: predict sentiment method in Multitask BERT class

```

def predict_similarity(self,
                        input_ids_1, attention_mask_1,
                        input_ids_2, attention_mask_2, noise1=0,
noise2=0, pooling='mean'):
    '''Given a batch of pairs of sentences, outputs a single logit
    corresponding to how similar they are.

```

```

    Note that your output should be unnormalized (a logit).
    '''

    ### TODO
    embedding_output1=self.forward(input_ids_1, attention_mask_1,
    noise1)
    embedding_output2=self.forward(input_ids_2, attention_mask_2,
    noise2)
    if pooling == 'mean':
        hidden_states1=embedding_output1['last_hidden_state']
        hidden_states2=embedding_output2['last_hidden_state']
        mask1=attention_mask_1.unsqueeze(-1).expand(-1, -1,
hidden_states1.shape[2]).float()
        mask2=attention_mask_2.unsqueeze(-1).expand(-1, -1,
hidden_states1.shape[2]).float()
        den1=torch.sum(attention_mask_1,dim=1).unsqueeze(-1).
expand(-1, hidden_states1.shape[2])+1e-9
        den2=torch.sum(attention_mask_2,dim=1).unsqueeze(-1).
expand(-1, hidden_states1.shape[2])+1e-9

        vector1 = torch.sum(mask1*(embedding_output1['
last_hidden_state']), dim=1)/den1
        vector2 = torch.sum(mask2*(embedding_output2['
last_hidden_state']), dim=1)/den2
    if pooling=='cls':
        vector1=embedding_output1['pooler_output']
        vector2=embedding_output2['pooler_output']
        Dropout_layer_output1=self.dropout_layer(vector1)
        Dropout_layer_output2=self.dropout_layer(vector2)
        Cosine_similarity=(torch.sum(Dropout_layer_output1*
Dropout_layer_output2,dim=1)/((torch.norm(Dropout_layer_output1
,dim=1)*torch.norm(Dropout_layer_output2,dim=1)))+1+1e-8)*2.5
    if isinstance(noise1, torch.Tensor) and isinstance(noise2,
torch.Tensor):
        if pooling=='mean':
            vector1_wn = torch.sum(mask1*(embedding_output1['
last_hidden_state_wn']), dim=1)/den1
            vector2_wn = torch.sum(mask2*(embedding_output2['
last_hidden_state_wn']), dim=1)/den2
        if pooling=='cls':
            vector1_wn=embedding_output1['pooler_output_wn']
            vector2_wn=embedding_output2['pooler_output_wn']
            Dropout_layer_output1_wn=self.dropout_layer(vector1_wn)
            Dropout_layer_output2_wn=self.dropout_layer(vector2_wn)
            Cosine_similarity_wn=(torch.sum(Dropout_layer_output1_wn*
Dropout_layer_output2_wn,dim=1)/((torch.norm(
Dropout_layer_output1_wn,dim=1)*torch.norm(
Dropout_layer_output2_wn,dim=1)))+1+1e-8)*2.5

```

```

        Embedding_in1=embedding_output1['embedding_output']
        Embedding_in2=embedding_output2['embedding_output']
        Embedding_in1_wn=embedding_output1['
embedding_output_with_noise']
        Embedding_in2_wn=embedding_output2['
embedding_output_with_noise']
        return Cosine_similarity, Cosine_similarity_wn,
Embedding_in1, Embedding_in2, Embedding_in1_wn, Embedding_in2_wn
    return Cosine_similarity

```

LISTING A.6: predict similarity method in Multitask BERT class

```

def predict_paraphrase(self,
                        input_ids_1, attention_mask_1,
                        input_ids_2, attention_mask_2, noise1=0,
noise2=0, pooling='mean'):
    '''Given a batch of pairs of sentences, outputs a single logit
    for predicting whether they are paraphrases.
    Note that your output should be unnormalized (a logit); it
    will be passed to the sigmoid function
    during evaluation.
    '''
    ### TODO
    embedding_output1=self.forward(input_ids_1, attention_mask_1,
noise1)
    embedding_output2=self.forward(input_ids_2, attention_mask_2,
noise2)
    if pooling=='mean':
        hidden_states1=embedding_output1['last_hidden_state']
        hidden_states2=embedding_output2['last_hidden_state']
        mask1=attention_mask_1.unsqueeze(-1).expand(-1, -1,
hidden_states1.shape[2]).float()
        mask2=attention_mask_2.unsqueeze(-1).expand(-1, -1,
hidden_states1.shape[2]).float()
        den1=torch.sum(attention_mask_1,dim=1).unsqueeze(-1).
expand(-1, hidden_states1.shape[2])+1e-9
        den2=torch.sum(attention_mask_2,dim=1).unsqueeze(-1).
expand(-1, hidden_states1.shape[2])+1e-9

        vector1 = torch.sum(mask1*(embedding_output1['
last_hidden_state']), dim=1)/den1
        vector2 = torch.sum(mask2*(embedding_output2['
last_hidden_state']), dim=1)/den2
    elif pooling=='cls':
        vector1=embedding_output1['pooler_output']
        vector2=embedding_output2['pooler_output']
    Dropout_layer_output1=self.dropout_layer(vector1)

```

```

Dropout_layer_output2=self.dropout_layer(vector2)

Concatenated_result=torch.concat((Dropout_layer_output1,
Dropout_layer_output2,abs(Dropout_layer_output1-
Dropout_layer_output2)),dim=1)
logit=self.paraphrase_classifier_layer(Concatenated_result)
if isinstance(noise1, torch.Tensor) and isinstance(noise2,
torch.Tensor):
    #mean
    if pooling=='mean':
        vector1_wn = torch.sum(mask1*(embedding_output1['
last_hidden_state_wn']), dim=1)/den1
        vector2_wn = torch.sum(mask2*(embedding_output2['
last_hidden_state_wn']), dim=1)/den2
    if pooling=='cls':
        vector1_wn=embedding_output1['pooler_output_wn']
        vector2_wn=embedding_output2['pooler_output_wn']
        Dropout_layer_output1_wn=self.dropout_layer(vector1_wn)
        Dropout_layer_output2_wn=self.dropout_layer(vector2_wn)
        Concatenated_result=torch.concat((Dropout_layer_output1_wn
,Dropout_layer_output2_wn,abs(Dropout_layer_output1_wn-
Dropout_layer_output2_wn)),dim=1)
        logit_wn=self.paraphrase_classifier_layer(
Concatenated_result)
        Embedding_in1=embedding_output1['embedding_output']
        Embedding_in2=embedding_output2['embedding_output']
        Embedding_in1_wn=embedding_output1['
embedding_output_with_noise']
        Embedding_in2_wn=embedding_output2['
embedding_output_with_noise']
        return logit,logit_wn,Embedding_in1,Embedding_in2,
Embedding_in1_wn,Embedding_in2_wn
    return logit

```

LISTING A.7: predict paraphrase method in Multitask BERT class

## Appendix B

# Implementation of Supervised SimCSE in Python (Abdul Samad)

### B.1 Supervised SimCSE

This section describes the implementation of Supervised SimCSE for contrastive learning using the SNLI dataset, including data preprocessing, triplet contrastive loss, and the training loop.

#### B.1.1 SNLI Dataset Preprocessing

To adapt the SNLI dataset for contrastive learning, we extract only the **entailment** (label = 0) and **contradiction** (label = 2) examples, ignoring **neutral** (label = 1) examples due to their lack of strong positive or negative signals. Each training instance is converted into a triplet: the premise as the anchor, the entailment hypothesis as the positive, and the contradiction hypothesis as the negative. Listing B.1 shows the preprocessing code.

```
import pandas as pd

# Load the SNLI CSV
df = pd.read_csv("snli-train.csv")

# Keep only entailment (0) and contradiction (2)
df = df[df['label'].isin([0, 2])]

# Separate positive and negative examples
positives = df[df['label'] == 0][['premise', 'hypothesis']].rename(
    columns={'hypothesis': 'positive_sentence'})
negatives = df[df['label'] == 2][['premise', 'hypothesis']].rename(
    columns={'hypothesis': 'negative_sentence'})
```

```

# Merge on 'premise' to form triplets
merged = pd.merge(positives, negatives, on='premise')

# Save triplet format
merged.to_csv("snli_processed.csv", index=False, encoding='utf-8')
print("New CSV file 'snli_processed.csv' created successfully.")

```

LISTING B.1: Preprocessing SNLI Dataset

### B.1.2 Triplet Contrastive Loss for SimCSE

The Supervised SimCSE objective maximizes the similarity between anchor and positive embeddings while minimizing it between anchor and negative embeddings. Given embeddings  $a_i$  (anchor: premise),  $p_i$  (positive: entailment hypothesis), and  $n_i$  (negative: contradiction hypothesis), all  $\ell_2$ -normalized, we compute cosine similarities:

$$s_i^+ = a_i^\top p_i, \quad s_i^- = a_i^\top n_i$$

A temperature-scaled softmax is applied:

$$P_i = \frac{\exp(s_i^+/\tau)}{\exp(s_i^+/\tau) + \exp(s_i^-/\tau)}$$

The instance-wise loss is:

$$\mathcal{L}_i = -\log P_i$$

The batch loss is the average:

$$\mathcal{L} = \frac{1}{N} \sum_{i=1}^N \mathcal{L}_i$$

This loss pulls  $p_i$  closer to  $a_i$  and pushes  $n_i$  away.

### B.1.3 Supervised SimCSE Loss Implementation

Listing B.2 implements the Supervised SimCSE loss within the MultitaskBERT class. It encodes the triplet inputs, applies mean pooling (or CLS pooling), normalizes embeddings, computes cosine similarities, and calculates the contrastive loss.

```

class MultitaskBERT(nn.Module):
    def supervised_simcse_loss(self, input_ids_1, attention_mask_1
        ,
                                input_ids_2, attention_mask_2,
                                input_ids_3, attention_mask_3,
                                pooling='mean', temperature=0.05):
        # Encode triplet input
        embedding_output1 = self.forward(input_ids_1,
        attention_mask_1)

```

```

        embedding_output2 = self.forward(input_ids_2,
        attention_mask_2)
        embedding_output3 = self.forward(input_ids_3,
        attention_mask_3)

        if pooling == 'mean':
            hidden_state1 = embedding_output1['last_hidden_state']
            hidden_state2 = embedding_output2['last_hidden_state']
            hidden_state3 = embedding_output3['last_hidden_state']

            mask1 = attention_mask_1.unsqueeze(-1).expand_as(
hidden_state1).float()
            mask2 = attention_mask_2.unsqueeze(-1).expand_as(
hidden_state2).float()
            mask3 = attention_mask_3.unsqueeze(-1).expand_as(
hidden_state3).float()

            denom1 = torch.sum(mask1, dim=1).unsqueeze(-1)
            denom2 = torch.sum(mask2, dim=1).unsqueeze(-1)
            denom3 = torch.sum(mask3, dim=1).unsqueeze(-1)

            vector1 = torch.sum(mask1 * hidden_state1, dim=1) /
denom1
            vector2 = torch.sum(mask2 * hidden_state2, dim=1) /
denom2
            vector3 = torch.sum(mask3 * hidden_state3, dim=1) /
denom3
        elif pooling == 'cls':
            vector1 = embedding_output1['pooler_output']
            vector2 = embedding_output2['pooler_output']
            vector3 = embedding_output3['pooler_output']

        anchor = self.dropout_layer(vector1)
        positive = self.dropout_layer(vector2)
        negative = self.dropout_layer(vector3)

        anchor = F.normalize(anchor, dim=-1)
        positive = F.normalize(positive, dim=-1)
        negative = F.normalize(negative, dim=-1)

        sim_ap = torch.matmul(anchor, positive.T) / temperature
        sim_an = torch.matmul(anchor, negative.T) / temperature

        exp_ap = torch.exp(torch.diag(sim_ap))
        denom = exp_ap + torch.exp(sim_an)

        loss = -torch.log(exp_ap / denom)

```

```
return loss.mean()
```

LISTING B.2: Supervised SimCSE Loss

### B.1.4 Training Loop for Supervised SimCSE

Listing B.3 shows the training loop for Supervised SimCSE using the SNLI triplet data. The model is trained with the AdamW optimizer, and performance is evaluated on both the SNLI development set and the STS dataset, saving the model when the STS development correlation improves.

```
optimizer = AdamW(model.parameters(), lr=args.lr)

for epoch in range(args.epochs):
    model.train()
    simcse_train_loss = 0
    num_batches = 0

    for batch in tqdm(snli_train_dataloader, desc=f'SimCSE-{epoch}'):
        b_ids1, b_mask1 = batch['token_ids_1'].to(device), batch['attention_mask_1'].to(device)
        b_ids2, b_mask2 = batch['token_ids_2'].to(device), batch['attention_mask_2'].to(device)
        b_ids3, b_mask3 = batch['token_ids_3'].to(device), batch['attention_mask_3'].to(device)

        optimizer.zero_grad()
        loss = model.supervised_simcse_loss(b_ids1, b_mask1, b_ids2, b_mask2, b_ids3, b_mask3)
        loss.backward()
        optimizer.step()

        simcse_train_loss += loss.item()
        num_batches += 1

    simcse_train_loss /= num_batches
    simcse_dev_acc = model_eval_simcse(snli_dev_dataloader, model, device)

    # Evaluate on STS
    dev_sts_corr = model_eval_sts(sts_dev_dataloader, model, device)
    if dev_sts_corr > best_sts_corr:
        best_sts_corr = dev_sts_corr
        save_model(model, optimizer, args, config_args, f"SimCSE-{epoch}")
```



```
print(f"Epoch {epoch}: simcse train loss :: {simcse_train_loss
      :.3f}, simcse dev acc :: {simcse_dev_acc :.3f}, dev sts corr
      :: {dev_sts_corr :.3f}")
```

LISTING B.3: Training SimCSE with Triplet SNLI Data

### B.1.5 SNLI Dataset Preparation

To prepare the SNLI dataset for training, we implemented custom dataset classes to handle the triplet structure and tokenization. Listing B.4 shows the `TrainSNLIDataset` and `SNLIDataset` classes. `TrainSNLIDataset` processes triplets (premise, positive, negative) for contrastive learning, tokenizing each sentence using the BERT tokenizer and returning padded token IDs, attention masks, and token type IDs for all three inputs. `SNLIDataset` handles paired data (e.g., for STS or QQP tasks), tokenizing two sentences and optionally converting labels to regression or classification format. Both classes include a `collate_fn` method to batch and pad the data for efficient training.

```
class TrainSNLIDataset(Dataset):
    def _init_(self, dataset, args):
        self.dataset = dataset
        self.p = args
        self.tokenizer = BertTokenizer.from_pretrained('bert-base-
uncased')

    def _len_(self):
        return len(self.dataset)

    def _getitem_(self, idx):
        return self.dataset[idx]

    def pad_data(self, data):
        sent1 = [x[0] for x in data]
        sent2 = [x[1] for x in data]
        sent3 = [x[2] for x in data]

        encoding1 = self.tokenizer(
            sent1,
            return_tensors='pt',
            padding=True,
            truncation=True
        )
        token_ids = torch.LongTensor(encoding1['input_ids'])
        attention_mask = torch.LongTensor(encoding1['
attention_mask'])
        token_type_ids = torch.LongTensor(encoding1['
token_type_ids'])
```

```

        encoding2 = self.tokenizer(
            sent2,
            return_tensors='pt',
            padding=True,
            truncation=True
        )
        token_ids2 = torch.LongTensor(encoding2['input_ids'])
        attention_mask2 = torch.LongTensor(encoding2['
attention_mask'])
        token_type_ids2 = torch.LongTensor(encoding2['
token_type_ids'])

        encoding3 = self.tokenizer(
            sent3,
            return_tensors='pt',
            padding=True,
            truncation=True
        )
        token_ids3 = torch.LongTensor(encoding3['input_ids'])
        attention_mask3 = torch.LongTensor(encoding3['
attention_mask'])
        token_type_ids3 = torch.LongTensor(encoding3['
token_type_ids'])

        return (token_ids, token_type_ids, attention_mask,
                token_ids2, token_type_ids2, attention_mask2,
                token_ids3, token_type_ids3, attention_mask3)

    def collate_fn(self, all_data):
        (token_ids, token_type_ids, attention_mask,
         token_ids2, token_type_ids2, attention_mask2,
         token_ids3, token_type_ids3, attention_mask3) = self.
pad_data(all_data)

        batched_data = {
            'token_ids_1': token_ids,
            'token_type_ids_1': token_type_ids,
            'attention_mask_1': attention_mask,
            'token_ids_2': token_ids2,
            'token_type_ids_2': token_type_ids2,
            'attention_mask_2': attention_mask2,
            'token_ids_3': token_ids3,
            'token_type_ids_3': token_type_ids3,
            'attention_mask_3': attention_mask3
        }

        return batched_data

```

```

class SNLIDataset(Dataset):
    def _init_(self, dataset, args, isRegression=False):
        self.dataset = dataset
        self.p = args
        self.isRegression = isRegression
        self.tokenizer = BertTokenizer.from_pretrained('bert-base-uncased')

    def _len_(self):
        return len(self.dataset)

    def _getitem_(self, idx):
        return self.dataset[idx]

    def pad_data(self, data):
        sent1 = [x[0] for x in data]
        sent2 = [x[1] for x in data]
        labels = [x[2] for x in data]

        encoding1 = self.tokenizer(sent1, return_tensors='pt', padding=True, truncation=True)
        encoding2 = self.tokenizer(sent2, return_tensors='pt', padding=True, truncation=True)

        token_ids = torch.LongTensor(encoding1['input_ids'])
        attention_mask = torch.LongTensor(encoding1['attention_mask'])
        token_type_ids = torch.LongTensor(encoding1['token_type_ids'])

        token_ids2 = torch.LongTensor(encoding2['input_ids'])
        attention_mask2 = torch.LongTensor(encoding2['attention_mask'])
        token_type_ids2 = torch.LongTensor(encoding2['token_type_ids'])

        if self.isRegression:
            labels = torch.DoubleTensor(labels)
        else:
            labels = torch.LongTensor(labels)

        return (token_ids, token_type_ids, attention_mask, token_ids2, token_type_ids2, attention_mask2, labels)

    def collate_fn(self, all_data):
        (token_ids, token_type_ids, attention_mask,

```

```
token_ids2, token_type_ids2, attention_mask2,
labels) = self.pad_data(all_data)

batched_data = {
    'token_ids_1': token_ids,
    'token_type_ids_1': token_type_ids,
    'attention_mask_1': attention_mask,
    'token_ids_2': token_ids2,
    'token_type_ids_2': token_type_ids2,
    'attention_mask_2': attention_mask2,
    'labels': labels
}
return batched_data
```

LISTING B.4: SNLI Dataset Classes

## Appendix C

# SBERT Implementation for Paraphrase Detection and Semantic Similarity (Hussnain Amjad)

### C.1 Overview

This appendix presents the Sentence-BERT (SBERT) implementation for paraphrase detection and semantic textual similarity (STS) tasks. SBERT enhances BERT by employing a siamese network to produce sentence embeddings optimized for sentence-level tasks. The code for the `predict_similarity` and `predict_paraphrase` methods is provided in Appendix A (Listings A.6 and A.7), integrated into the `MultitaskBERT` class. In Listing A.6, we did not use any linear layer; instead, we used cosine similarity for the STS task, as recommended for regression tasks in *Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks* by Reimers and Gurevych [9]. For paraphrase detection, the concatenation of sentence embeddings was performed, as suggested in the same paper. For both tasks, we set mean pooling as the default and used it during training. The attention mask was applied to obtain mean embeddings from the last layer of BERT, ensuring that embeddings corresponding to padding tokens were excluded from the mean embedding calculation. The cosine similarity for Semantic Textual Similarity Task is scaled to be between 0 and 5.

# References

- [1] Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. Bert: Pre-training of deep bidirectional transformers for language understanding. *arXiv preprint arXiv:1810.04805*, 2018.
- [2] Tianyu Gao, Xingcheng Yao, and Danqi Chen. Simcse: Simple contrastive learning of sentence embeddings. *arXiv preprint arXiv:2104.08821*, 2021.
- [3] Haoming Jiang, Di He, Aniruddha Chen, Qipeng Guo Wang, Ming Song, and Richard Zemel. Smart: Robust and efficient fine-tuning for pre-trained natural language models through principled regularized optimization. *arXiv preprint arXiv:1911.03437*, 2019.
- [4] Akhil Mahankali et al. Cs224n spring 2024 dfp student handout. <https://github.com/amahankali10/CS224N-Spring2024-DFP-Student-Handout>, 2024. GitHub repository, accessed: 2025-05-29.
- [5] Tomas Mikolov, Kai Chen, Greg Corrado, and Jeffrey Dean. Efficient estimation of word representations in vector space. In *International Conference on Learning Representations (ICLR)*, 2013.
- [6] Papers with Code. Papers with code, 2024. URL <https://www.paperswithcode.com>. Accessed: 2025-05-12.
- [7] Jeffrey Pennington, Richard Socher, and Christopher Manning. Glove: Global vectors for word representation. In *Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 1532–1543, 2014.
- [8] Pranav Rajpurkar, Jian Zhang, Konstantin Lopyrev, and Percy Liang. Squad: 100,000+ questions for machine comprehension of text. In *Proceedings of the 2016 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 2383–2392, 2016.
- [9] Nils Reimers and Iryna Gurevych. Sentence-bert: Sentence embeddings using siamese bert-networks. *arXiv preprint arXiv:1908.10084*, 2019.
- [10] Stanford University. CS224n: Natural Language Processing with Deep Learning, 2024. URL <https://web.stanford.edu/class/archive/cs/cs224n/cs224n.1246/>. Accessed: 2025-05-29.

- 
- [11] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Lukasz Kaiser, and Illia Polosukhin. Attention is all you need. In *Advances in Neural Information Processing Systems*, volume 30. Curran Associates, Inc., 2017.